



AWS
re:Invent

ANT 418

Deep dive and best practices for Amazon Redshift


Tony Gibbs

Sr. Database Specialist SA
Amazon Web Services

Harshida Patel

Data Warehouse Specialist SA
Amazon Web Services

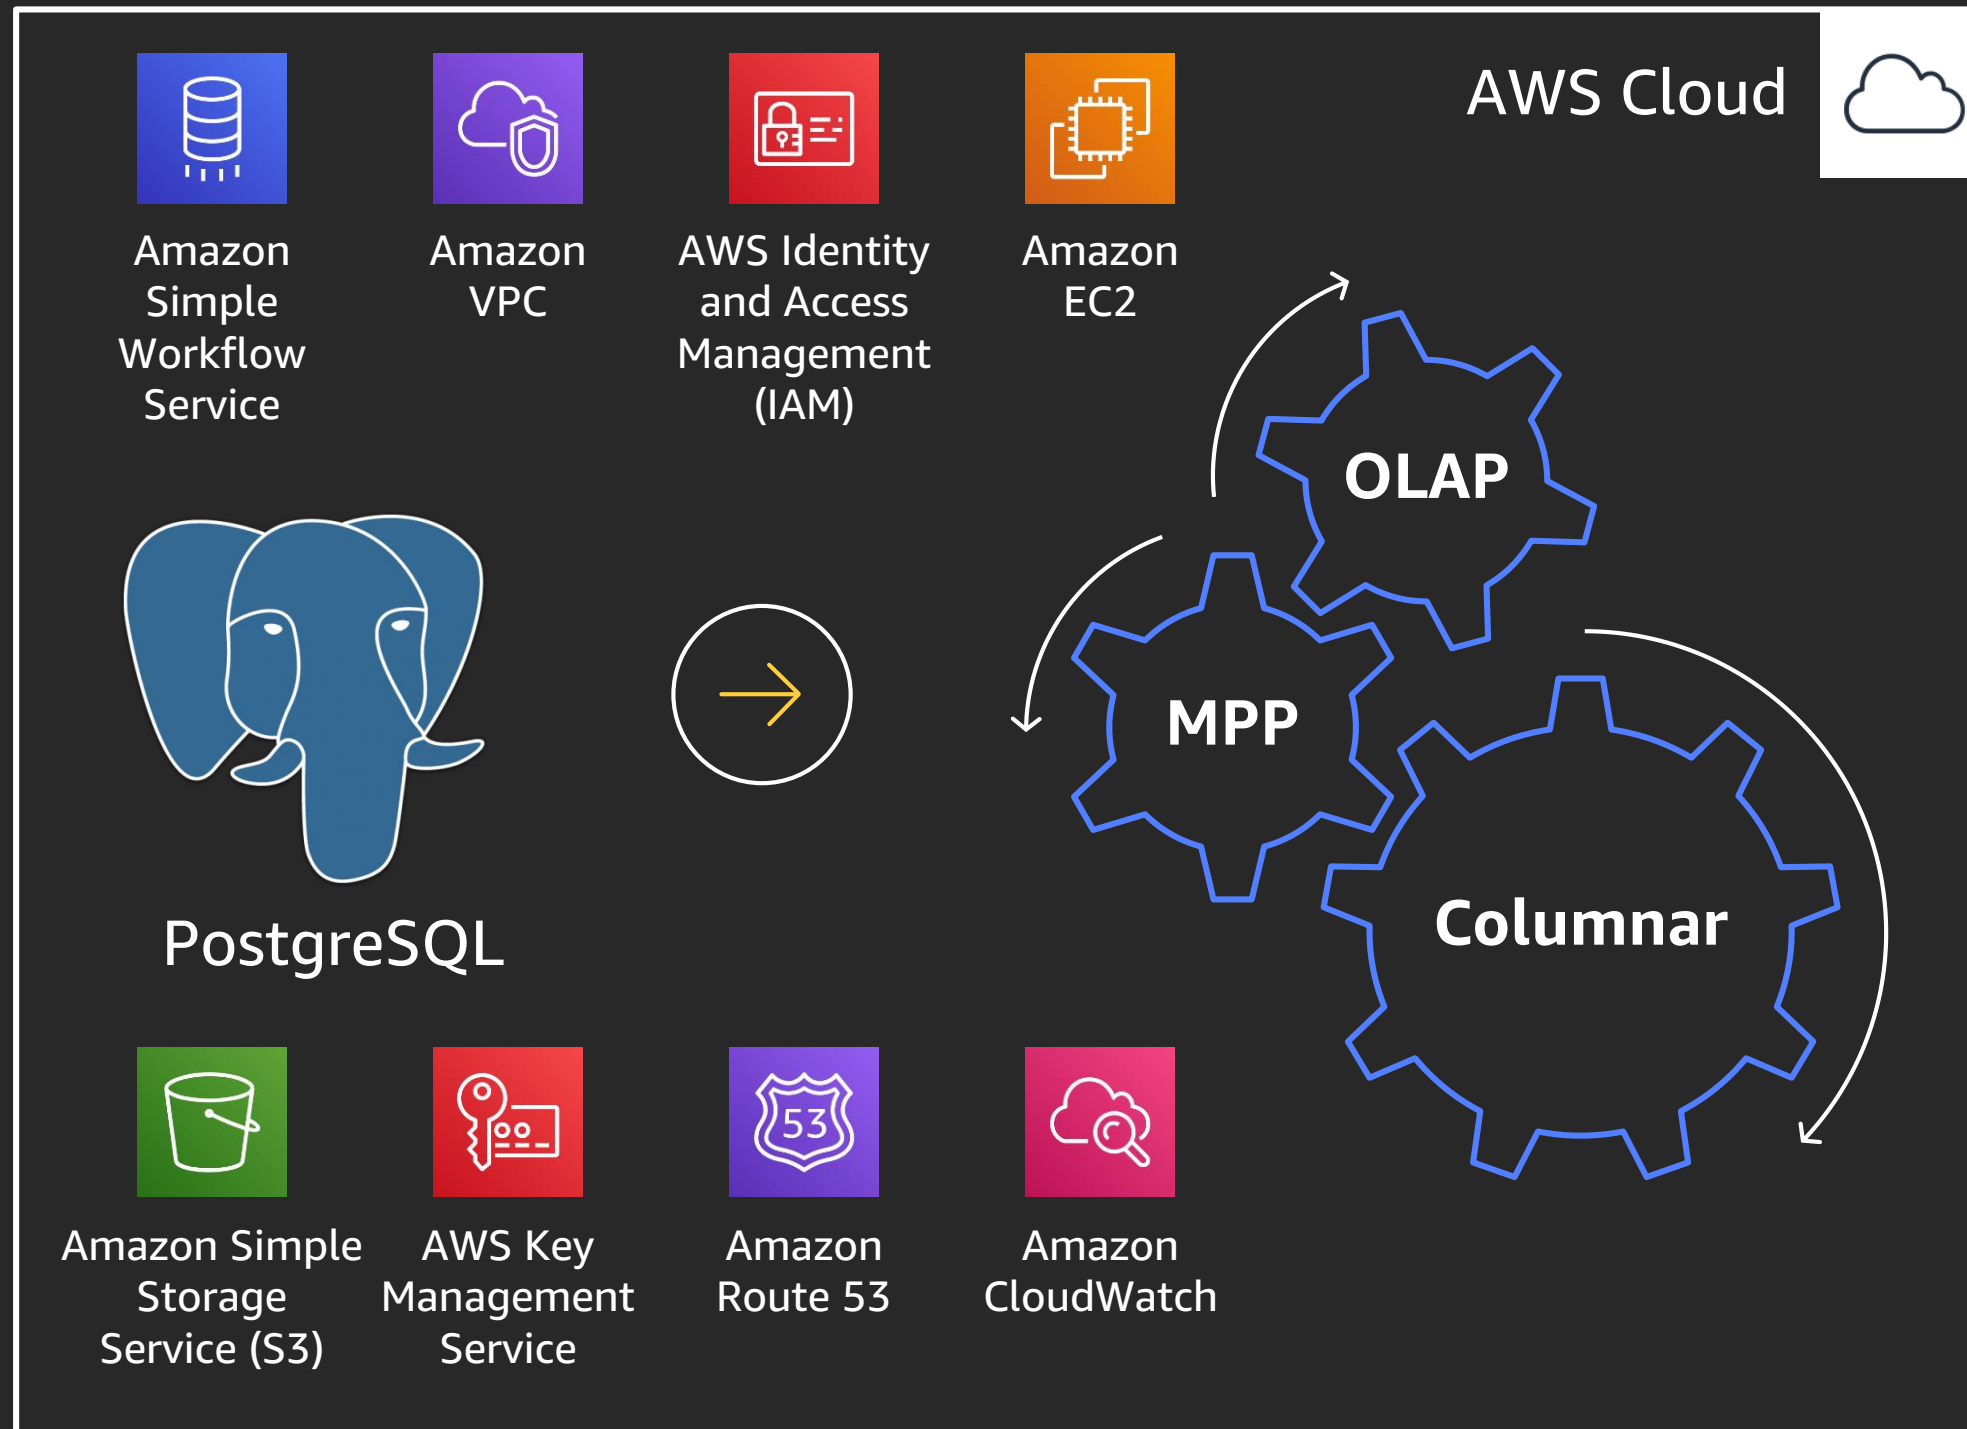
Are you an Amazon Redshift user?




Agenda

- Architecture and concepts
- Data storage, ingestion, and ELT
- Workload management and query monitoring rules
- Cluster sizing and resizing
- Amazon Redshift Advisor
- Additional resources
- Open Q&A


Architecture and concepts




February 2013



>175 significant patches



December 2019



Amazon Redshift has been innovating quickly

Robust result set caching

Amazon Redshift Spectrum: Date formats, scalar json and ION file formats support, region expansion, predicate filtering

Unload to CSV

Auto WLM

Concurrency Scaling

Manage multi-part query in AWS console

Redshift Spectrum: Row group filtering in Parquet and ORC, nested data support, enhanced VPC routing, multiple partitions

Auto WLM with query priorities

Spatial processing

Large # of tables support ~20000

Auto analyze

~25 Query Monitoring Rules (QMR) support

Auto analyze for incremental changes on table

Faster Classic resize with optimized data transfer protocol

Snapshot scheduler

Stored procedures

Column level access control with AWS Lake Formation

Copy command support for ORC, Parquet

Health and performance monitoring w/Amazon CloudWatch

Performance: Join pushdowns to subquery, mixed workloads temporary tables, rank functions, null handling in join, single row insert

RA3

IAM role chaining

Automatic table distribution style

Advisor recommendations for distribution keys

Performance of inter-region snapshot transfers

Federated query

Elastic resize

Amazon CloudWatch support for WLM queues

Spectrum Request Accelerator

Performance: Bloom filters in joins, complex queries that create internal table, communication layer

AWS Lake Formation integration

AZ64 compression encoding

Materialized views

Groups

Performance enhancements: Hash join, vacuum, window functions, resize ops, aggregations, console, union all, efficient compile code cache

AQUA

Resiliency of ROLLBACK processing

Apply new distribution key

Amazon Redshift Spectrum: Concurrency scaling

Auto-Vacuum sort, Auto-Analyze, and Auto Table Sort

Console redesign

Manual pause and resume

200+

new features and enhancements in the past 18 months

Amazon Redshift architecture

Massively parallel, shared nothing columnar architecture

Leader node

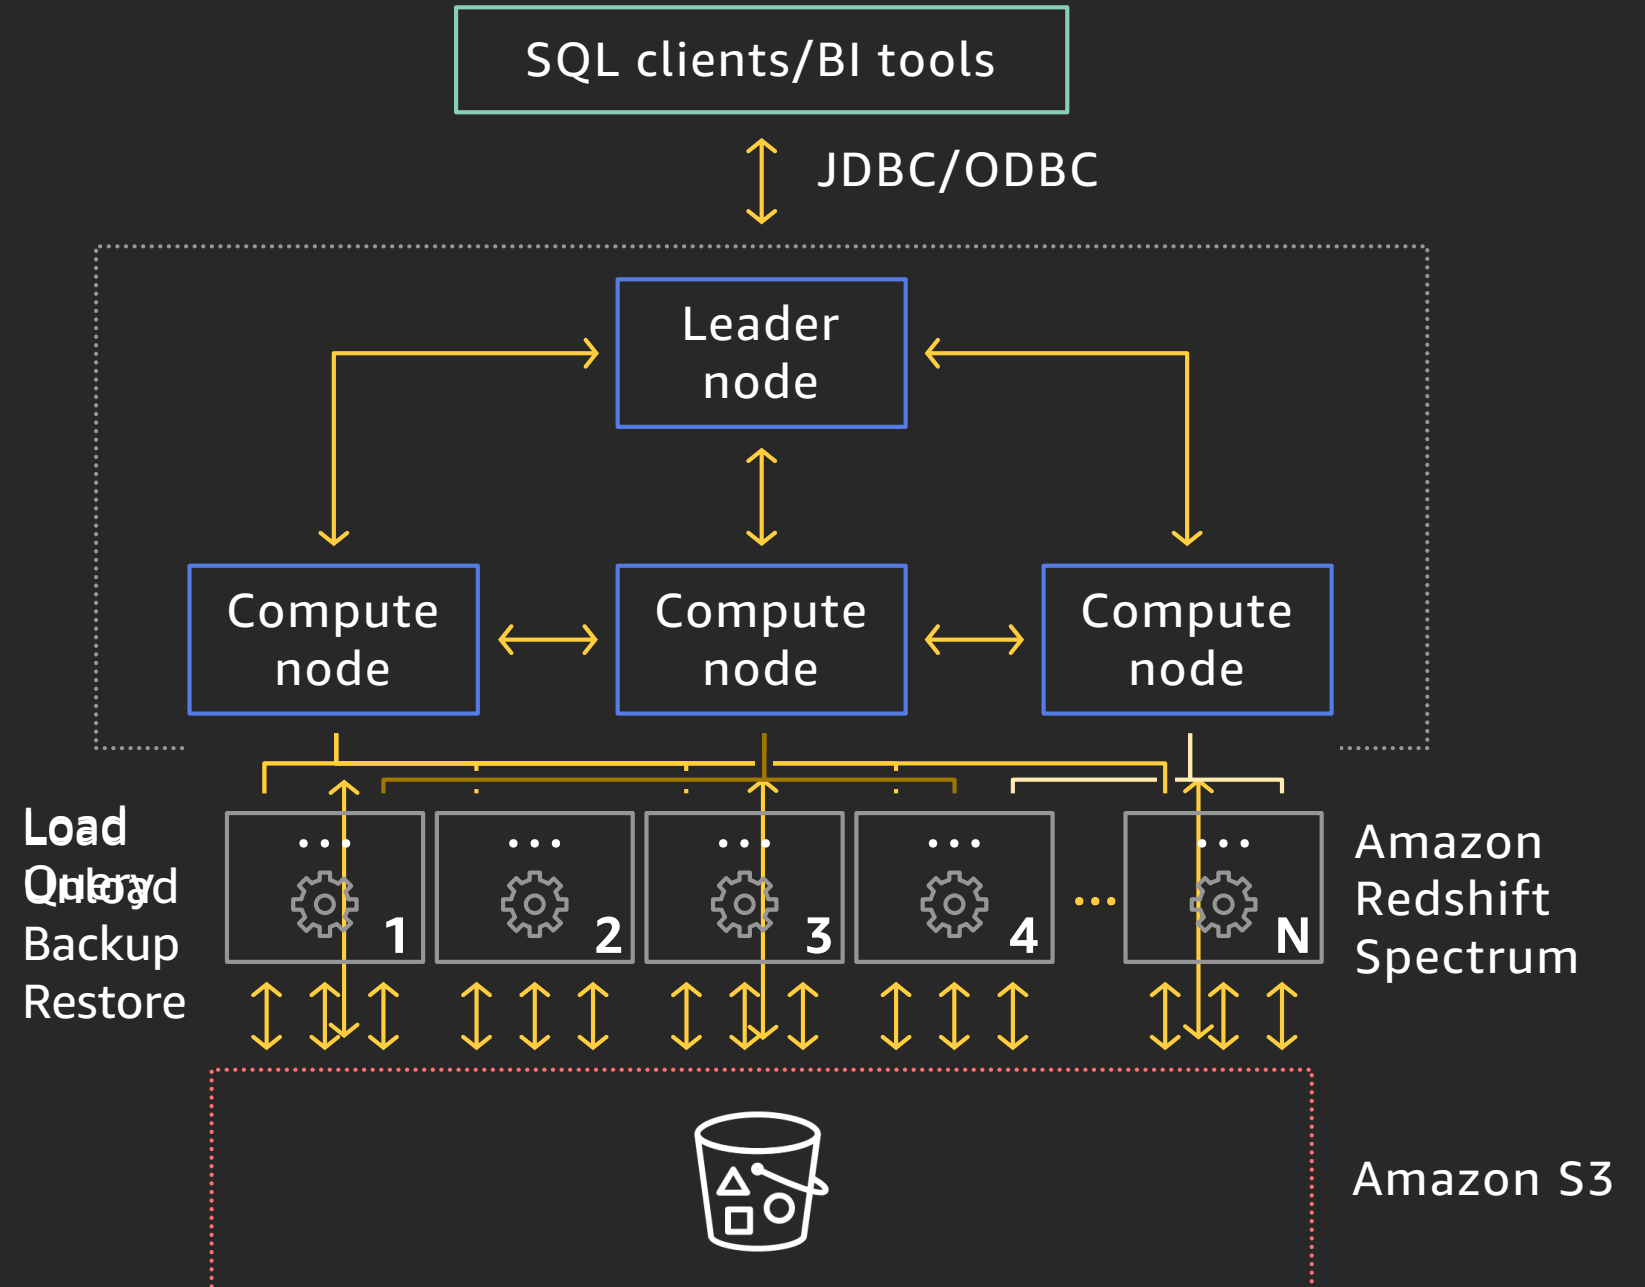
- SQL endpoint
- Stores metadata
- Coordinates parallel SQL processing

Compute nodes

- Local, columnar storage
- Executes queries in parallel
- Load, unload, backup, restore

Amazon Redshift Spectrum nodes

- Execute queries directly against Amazon Simple Storage Service (Amazon S3)



Amazon Redshift evolving architecture

Massively parallel, shared nothing
columnar architecture

Leader node

Compute nodes

Amazon Redshift Spectrum nodes

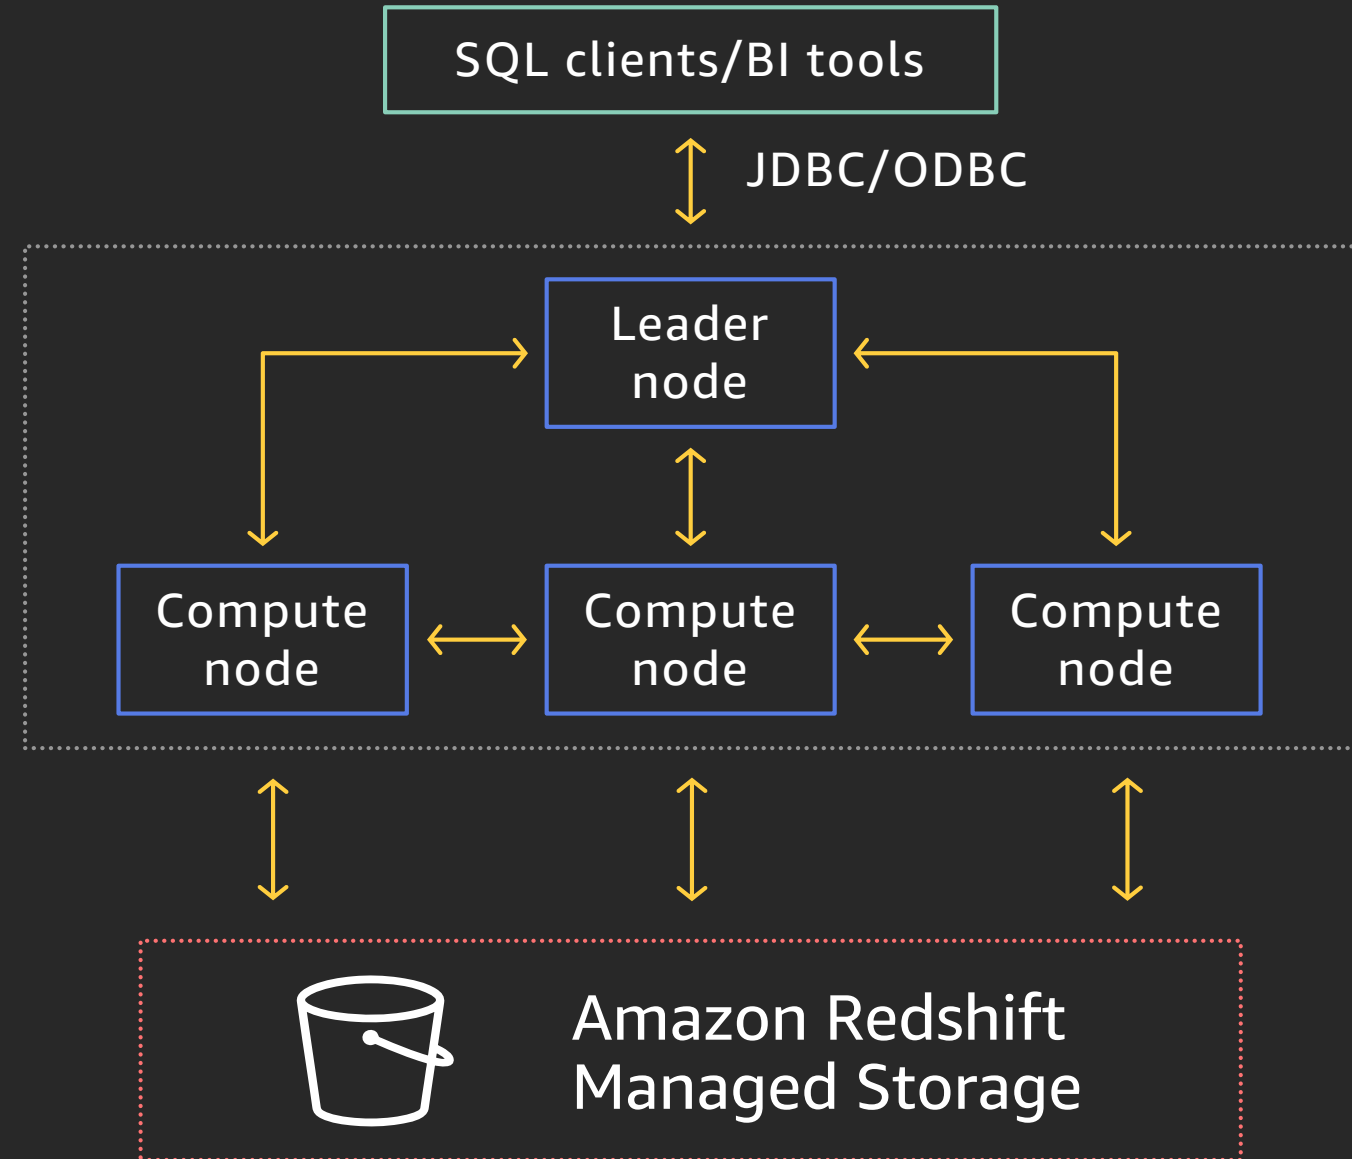
Amazon Redshift Managed Storage

Pay separately for storage and compute


Large high-speed SSD backed cache

Automatic scaling (up to 64 TB/instance)

Supports up to 8.2 PB of cluster storage



AQUA (Advanced Query Accelerator) for Amazon Redshift



New distributed & hardware-accelerated processing layer

With AQUA, Amazon Redshift is up to **10x faster** than any other cloud data warehouse, no extra cost

AQUA Nodes with custom AWS-designed analytics processors to make operations (compression, encryption, filtering, and aggregations) faster than traditional CPUs

Available in Preview with RA3. No code changes required

Terminology and concepts: Node types

Amazon Redshift analytics—RA3 (new)

Amazon Redshift Managed Storage (RMS)—Solid-state disks + Amazon S3

Dense compute—DC2

Solid-state disks

Dense storage—DS2

Magnetic disks

Instance type	Disk type	Size	Memory	CPU	Slices
RA3 4xlarge (coming soon)	RMS	Scales to 16 TB	96 GB	12	4
RA3 16xlarge (new)	RMS	Scales to 64 TB	384 GB	48	16
DC2 large	SSD	160 GB	16 GB	2	2
DC2 8xlarge	SSD	2.56 TB	244 GB	32	16
DS2 xlarge	Magnetic	2 TB	32 GB	4	2
DS2 8xlarge	Magnetic	16 TB	244 GB	36	16

Terminology and concepts: Columnar

Amazon Redshift uses a columnar architecture for storing data on disk

Goal: reduce I/O for analytics queries

Physically store data on disk by column rather than row

Only read the column data that is required

Columnar architecture: Example

```
CREATE TABLE deep_dive (  
  aid INT      --audience_id  
  ,loc CHAR(3) --location  
  ,dt  DATE    --date  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

```
SELECT min(dt) FROM deep_dive;
```

Row-based storage

- Need to read everything
- Unnecessary I/O

Columnar architecture: Example

```
CREATE TABLE deep_dive (  
  aid INT      --audience_id  
  ,loc CHAR(3) --location  
  ,dt  DATE    --date  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

```
SELECT min(dt) FROM deep_dive;
```

Column-based storage

- Only scan blocks for relevant column

Terminology and concepts: Compression

Goals

Allow more data to be stored within an Amazon Redshift cluster

Improve query performance by decreasing I/O

Impact

Allows two to four times more data to be stored within the cluster

By default, COPY automatically analyzes and compresses data on first load into an empty table

ANALYZE COMPRESSION is a built-in command that will find the optimal compression for each column on an existing table

Compression: Example

```
CREATE TABLE deep_dive (  
  aid INT      --audience_id  
  ,loc CHAR(3) --location  
  ,dt  DATE    --date  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

Add 1 of 13 different encodings
to each column

Compression: Example

```
CREATE TABLE deep_dive (  
  aid INT          ENCODE AZ64  
  ,loc CHAR(3)     ENCODE BYTEDICT  
  ,dt  DATE        ENCODE RUNLENGTH  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

More efficient compression is due to storing the same data type in the columnar architecture

Columns grow and shrink independently

Reduces storage requirements

Reduces I/O

New Amazon Redshift encoding type: AZ64

AZ64 is Amazon's proprietary compression encoding algorithm designed to achieve a high compression ratio and improved query processing

Goals:

Increase compression ratio, reducing the required footprint

Increase query performance by decreasing both encoding/decoding times

Result:

	AZ64 storage savings	AZ64 performance speed ups
RAW	60–70% less storage	25–30% faster
LZO	35% less storage	40% faster
ZSTD	Comparable footprint	70% faster

30TB [Cloud DW benchmark](#) is based on [TPC-DS \(v2.10\)](#) with no query modifications done

Best practices: Compression

Apply compression to all tables

In most cases, use AZ64 for INT, SMALLINT, BIGINT, TIMESTAMP, TIMESTAMPTZ, DATE, NUMERIC

In most cases, use LZO/ZSTD for VARCHAR and CHAR

Use ANALYZE COMPRESSION command to find optimal compression

RAW (no compression) for sparse columns and small tables

Changing column encoding requires a table rebuild

<https://github.com/awslabs/amazon-redshift-utils/tree/master/src/ColumnEncodingUtility>

Verifying columns are compressed:

```
SELECT "column", type, encoding FROM pg_table_def
WHERE tablename = 'deep_dive';
```

column	type	encoding
aid	integer	az64
loc	character(3)	bytedict
dt	date	runlength

Terminology and concepts: Blocks

Column data is persisted to 1 MB immutable blocks

Blocks are individually encoded with 1 of 13 encodings

A full block can contain millions of values

Terminology and concepts: Zone maps

Goal

Eliminates unnecessary I/O

In-memory block metadata

- Contains per-block min and max values

- All blocks automatically have zone maps

- Effectively prunes blocks that cannot contain data for a given query

Terminology and concepts: Data sorting

Goal

Make queries run faster by increasing the effectiveness of zone maps and reducing I/O

Impact

Enables range-restricted scans to prune blocks by leveraging zone maps

Achieved with the table property
SORTKEY defined on one or more columns

Optimal sort key is dependent on:

- Query patterns

- Business requirements

- Data profile

Sort key: Example

```
CREATE TABLE deep_dive (  
  aid INT      --audience_id  
  ,loc CHAR(3) --location  
  ,dt DATE     --date  
) SORTKEY (dt, loc);
```

Add a sort key to one or more columns to physically sort the data on disk

deep_dive		
aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14






deep_dive (sorted)		
aid	loc	dt
3	SFO	2017-04-01
4	JFK	2017-05-14
2	JFK	2017-10-20
1	SFO	2017-10-20


Zone maps and sorting: Example

```
SELECT count(*) FROM deep_dive WHERE dt = '06-09-2017';
```

Unsorted table

	MIN: 01-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 08-JUNE-2017 MAX: 30-JUNE-2017
	MIN: 12-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 02-JUNE-2017 MAX: 25-JUNE-2017

Sorted by date

	MIN: 01-JUNE-2017 MAX: 06-JUNE-2017
	MIN: 07-JUNE-2017 MAX: 12-JUNE-2017
	MIN: 13-JUNE-2017 MAX: 21-JUNE-2017
	MIN: 21-JUNE-2017 MAX: 30-JUNE-2017

Best practices: Sort keys

Place the sort key on columns that are frequently filtered on placing the lowest cardinality columns first

On most fact tables, the first sort key column should be a temporal column

Columns added to a sort key after a high-cardinality column are not effective

With an established workload, use the following scripts to help find sort key suggestions:

https://github.com/awslabs/amazon-redshift-utils/blob/master/src/AdminScripts/filter_used.sql

https://github.com/awslabs/amazon-redshift-utils/blob/master/src/AdminScripts/predicate_columns.sql

Design considerations:

Sort keys are less beneficial on small tables

Define four or less sort key columns—more will result in marginal gains and increased ingestion overhead

Terminology and concepts: Materialize columns

Goal: Make queries run faster by leveraging zonemaps on the fact tables

Frequently filtered and unchanging dimension values should be materialized within fact tables

Time dimension tables do not allow for range restricted scans on fact tables

Materializing temporal values in fact table can give significant performance gains

Example:

```
SELECT COUNT(*) FROM fact_dd JOIN dim_dd USING (timeid) WHERE dim_dd.ts >= '2018-11-29';  
SELECT COUNT(*) FROM fact_dd WHERE fact_dd.timestamp >= '2018-11-29';           -- Faster
```

Often calculated values should be materialized within fact tables

Example:

```
SELECT COUNT(*) FROM dd WHERE EXTRACT(EPOCH FROM ts) BETWEEN 1541120959 AND 1543520959;  
SELECT COUNT(*) FROM dd WHERE sorted_epoch BETWEEN 1541120959 AND 1543520959; -- Faster
```

Terminology and concepts: Slices

A slice can be thought of like a virtual compute node

Unit of data partitioning
Parallel query processing

Facts about slices

Each compute node is initialized with **either 2 or 16 slices**

Table rows are distributed to slices

A slice processes only its own data

Data distribution

Distribution style is a table property which dictates how that table's data is distributed throughout the cluster

KEY: Value is hashed, same value goes to same location (slice)

ALL: Full table data goes to the first slice of every node




EVEN: Round robin

AUTO: Combines EVEN and ALL

Goals

Distribute data evenly for parallel processing

Minimize data movement during query processing



Data distribution: Example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) (EVEN|KEY|ALL|AUTO);
```

Slice 0

Slice 1

Node 1

Table: deep_dive

User columns

System columns

aid

loc

dt

ins

del

row

Slice 2

Slice 3

Node 2

Data distribution: **EVEN** Example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE EVEN;
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 0					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 1					

Node 1

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 2					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 3					

Node 2

Data distribution: **KEY** Example #1

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (loc);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					

Slice 0

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					

Slice 1

Node 1

Rows: 0					
---------	--	--	--	--	--

Slice 2

Rows: 0					
---------	--	--	--	--	--

Slice 3

Node 2

Data distribution: **KEY** Example #2

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (aid);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 0					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 1					

Node 1

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 2					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 0					
Slice 3					

Node 2

Data distribution: **ALL** Example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE ALL;
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 

Slice 0

Rows: 0

Slice 1

Node 1

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 

Slice 2

Rows: 0

Slice 3

Node 2

Summary: Data distribution

DISTSTYLE **KEY**

Goals

- Optimize **JOIN** performance between large tables by distributing on columns used in the **ON** clause
- Optimize **INSERT INTO SELECT** performance
- Optimize **GROUP BY** performance

The column that is being distributed on should have a high cardinality and not cause row skew:

DISTSTYLE **ALL**

Goals

- Optimize **JOIN** performance with dimension tables
- Reduces disk usage on small tables

Small and medium size dimension tables (<3M rows)

DISTSTYLE **EVEN**

If neither **KEY** or **ALL** apply

DISTSTYLE **AUTO**

Default distribution—combines DISTSTYLE **ALL** and **EVEN**

```
SELECT diststyle, skew_rows  
FROM svv_table_info WHERE "table" = 'deep_dive';
```

diststyle	skew_rows
KEY(aid)	1.07

← Ratio between the slice with the most and least number of rows

Best practices: Table design summary

Add compression to columns

Use AZ64 where possible, ZSTD/LZO for most (VAR)CHAR columns

Add sort keys on the columns that are frequently filtered on

Materialize often filtered columns from dimension tables into fact tables

Materialize often calculated values into tables

Co-locate large tables using **DISTSTYLE KEY** if the columns do not cause skew

Avoid distribution keys on temporal columns

Keep data types as wide as necessary (but no longer than necessary)

VARCHAR, CHAR, and NUMERIC

Data storage, ingestion, and ELT

Terminology and concepts: Redundancy

Amazon Redshift DC/DS instances utilize locally attached storage devices
Amazon Redshift RA3 instances utilize Amazon Redshift Managed Storage

Global **commit** ensures all permanent tables have blocks written to multiple locations to ensure data redundancy

Asynchronously back up blocks to Amazon S3—in all cases, snapshots are transitionally consistent

Snapshot generated every 5 GB of changed data or eight hours

User can create on-demand manual snapshots

To disable backups at the table level: `CREATE TABLE example(id int) BACKUP NO;`

Temporary tables

Blocks are not mirrored to the remote partition—two-times faster write performance

Do not trigger a full commit or backups

Terminology and concepts: Transactions

Amazon Redshift is a fully transactional,
ACID-compliant data warehouse

Isolation level is serializable

Two-phase commits (local and global commit phases)

Cluster commit statistics

https://github.com/awslabs/amazon-redshift-utils/blob/master/src/AdminScripts/commit_stats.sql

Design consideration

Because of the expense of commit overhead, limit commits
by explicitly creating transactions

Data ingestion: COPY statement

Ingestion throughput

Each slice's query processors can load one file at a time:

Streaming decompression

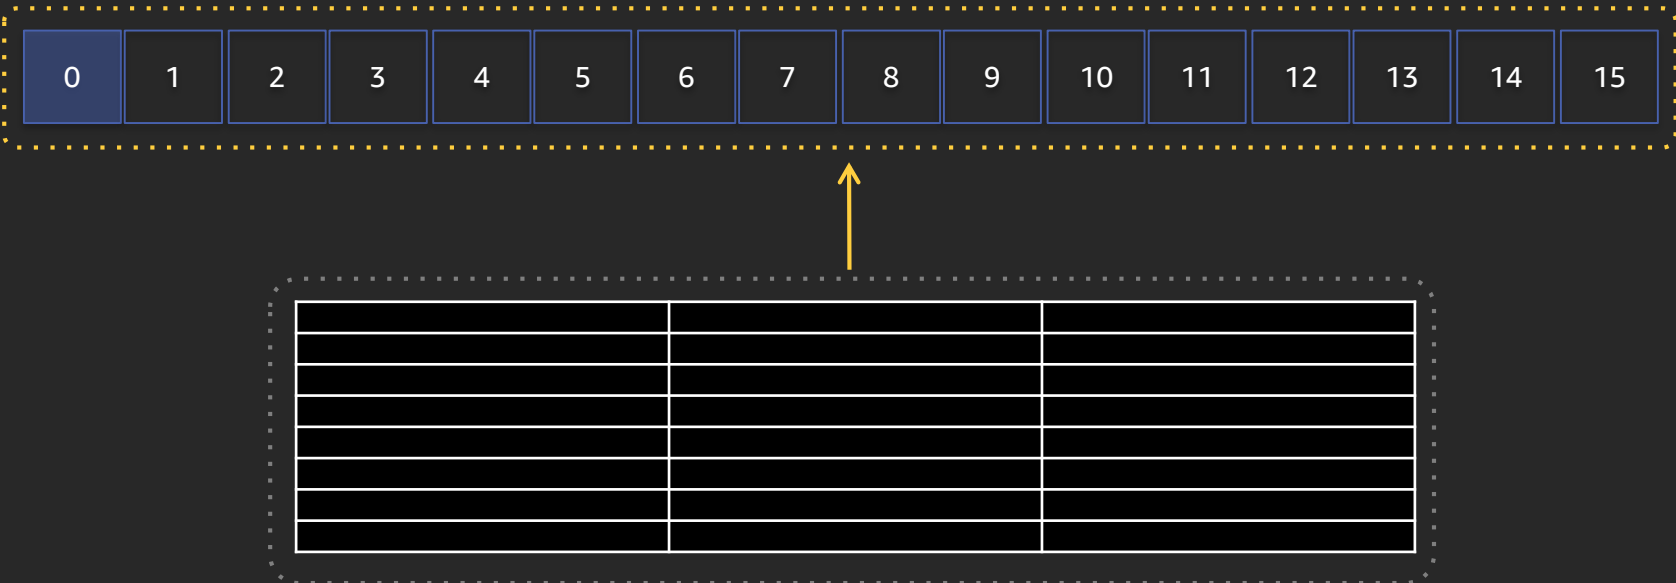
Parse

Distribute

Write

Realizing only partial node usage as 6.25% of slices are active

RA3.16XL compute node



1 input file

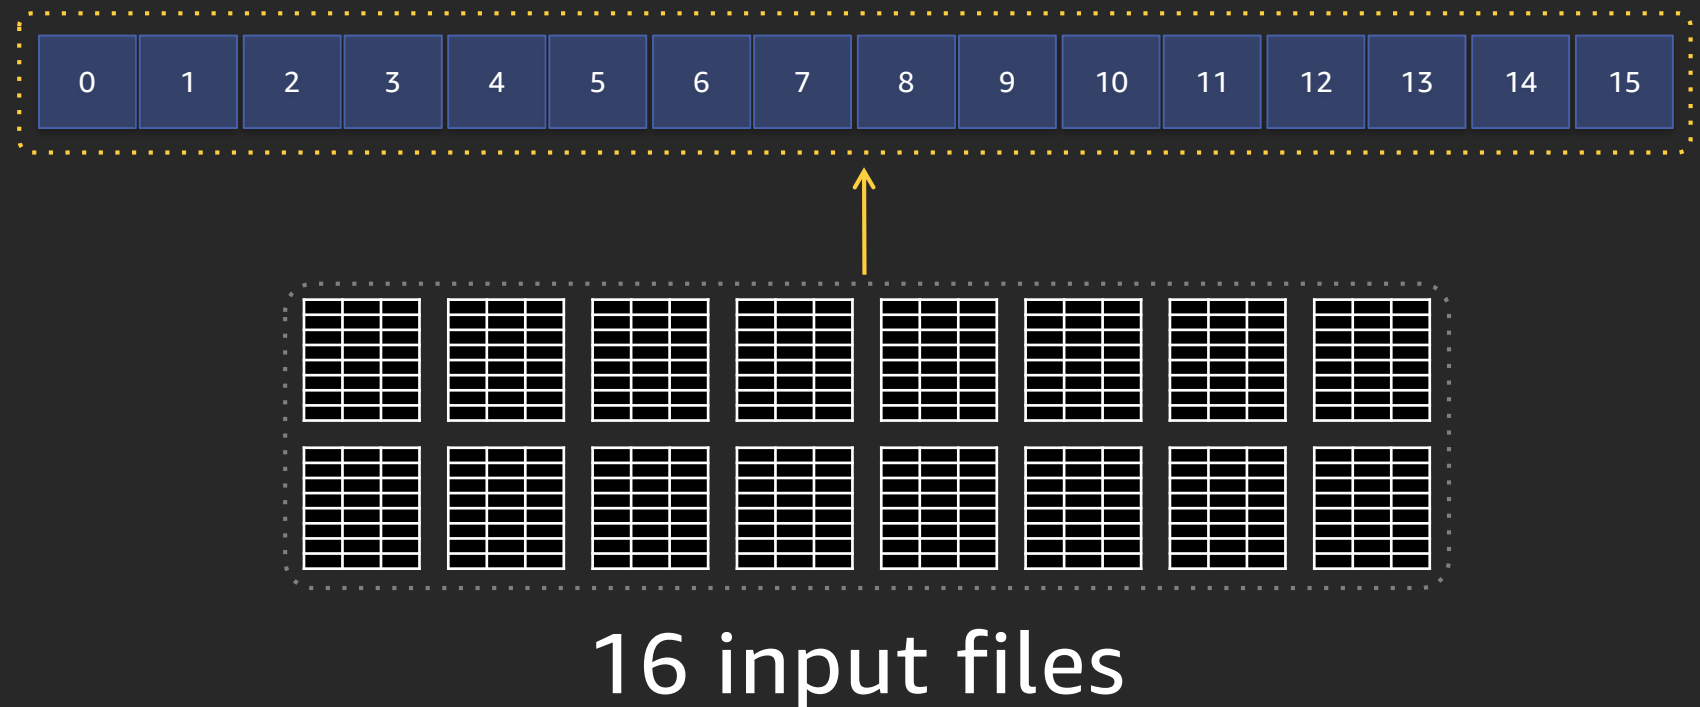
Data ingestion: COPY statement

Number of input files should be a multiple of the number of slices

Splitting the single file into 16 input files, all slices are working to maximize ingestion performance

COPY continues to scale linearly as you add nodes

RA3.16XL compute node



Recommendation is to use delimited files—1 MB to 1 GB after gzip compression

Best practices: COPY ingestion

Delimited files are recommended

Pick a simple delimiter '|' or ',' or tabs

Pick a simple NULL character (\N)

Use double quotes and an escape character (' \ ') for varchars

UTF-8 varchar columns take four bytes per char

Split files into a number that is a multiple of the total number of slices in the Amazon Redshift cluster

```
SELECT count(slice) from stv_slices;
```

Files sizes should be 1 MB to 1 GB after gzip compression

Data ingestion: Amazon Redshift Spectrum

Use **INSERT INTO SELECT** from external Amazon S3 tables

Aggregate incoming data

Select subset of columns and/or rows

Manipulate incoming column data with SQL

Load data in alternative file formats: Amazon ION, Grok, RCFile, and Sequence

Best practices

Save cluster resources for querying and reporting rather than on ELT

Filtering/aggregating incoming data can improve performance over COPY

Design considerations

Repeated reads against Amazon S3 are not transactional

\$5/TB of (compressed) data scanned

Design considerations: Data ingestion

Designed for large writes

Batch processing system, optimized for processing massive amounts of data

1 MB size plus immutable blocks means that we clone blocks on write so as not to introduce fragmentation

Small write (~1–10 rows) has similar cost to a larger write (~100K rows)

UPDATE and DELETE

Immutable blocks means that we only logically delete rows on UPDATE or DELETE

(AUTO) VACUUM to remove ghost rows from table

Data ingestion: Deduplication/UPsert

Table: deep_dive		
aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14
5	SJC	2017-10-10
6	SEA	2017-11-29



s3://bucket/dd.csv		
aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
5	SJC	2017-10-10
6	SEA	2017-11-29

Data ingestion: Deduplication/UPsert

Steps:

1. Load CSV data into a staging table

2. Delete duplicate data from the production table

3. Insert (or append) data from the staging into the production table

Data ingestion: Deduplication/UPSER

BEGIN;

CREATE **TEMP** TABLE staging(**LIKE** deep_dive);

COPY staging FROM 's3://bucket/dd.csv'
: 'creds' **COMPUPDATE OFF**;

DELETE FROM deep_dive d

USING staging s **WHERE** d.aid = s.aid;

INSERT INTO deep_dive **SELECT** * **FROM** staging;

DROP TABLE staging;

COMMIT;

Best practices: ELT

Wrap workflow/statements in an explicit transaction

Consider using **DROP TABLE** or **TRUNCATE** instead of **DELETE**

Staging tables:

Use temporary table or permanent table with the "**BACKUP NO**" option

If possible, use **DISTSTYLE KEY** on both the staging table and production table to speed up the **INSERT INTO SELECT** statement

With **COPY**, turn off automatic compression—**COMPUPDATE OFF**

Copy compression settings from the production table (using **LIKE** keyword) or manually apply compression to **CREATE TABLE DDL** (from **ANALYZE COMPRESSION** output)

For copying a large number of rows (> hundreds of millions), consider using **ALTER TABLE APPEND** instead of **INSERT INTO SELECT**

(AUTO) VACUUM

The VACUUM process runs either manually or automatically in the background

Goals

VACUUM will remove rows that are marked as deleted

VACUUM will globally sort tables

For tables with a sort key, ingestion operations will locally sort new data and write it into the unsorted region

Best practices

VACUUM should be run only as necessary

For the majority of workloads, **AUTO VACUUM DELETE** will reclaim space and **AUTO TABLE SORT** will sort the needed portions of the table

In cases where you know your workload—VACUUM can be run manually

Use **VACUUM BOOST** at off peak times (blocks deletes), which is as quick as “Deep Copy”

(AUTO) ANALYZE

The ANALYZE process collects table statistics for optimal query planning

In the vast majority of cases, **AUTO ANALYZE automatically handles statistics gathering**

Best practices

ANALYZE can be run periodically after ingestion on just the columns that WHERE predicates are filtered on

Utility to manually run VACUUM and ANALYZE on all the tables in the cluster:
<https://github.com/awslabs/amazon-redshift-utils/tree/master/src/AnalyzeVacuumUtility>

Workload management and query monitoring rules

Workload management (WLM)

Allows for the separation of different query workloads

Goals

Prioritize important queries

Throttle/abort less important queries

Control concurrent number of executing queries

Divide cluster memory

Set query timeouts to abort long running queries

Terminology and concepts: WLM attributes

Queues

Assigned a percentage of cluster memory

SQL queries execute in queue based on

User group: which groups the user belongs to

Query group session level variable

Query slots (or Concurrency):

Division of memory within a WLM queue, correlated with the number of simultaneous running queries

WLM_QUERY_SLOT_COUNT is a session level variable

Useful to increase for memory intensive operations
(example: large COPY, VACUUM, large INSERT INTO SELECT)

Terminology and concepts: WLM attributes

Short query acceleration (SQA)


Automatically detect short running queries and run them within the short query queue if queuing occurs

Concurrency scaling

When queues are full, queries are routed to transient Amazon Redshift clusters

Concurrency scaling

Amazon Redshift automatically adds transient clusters, in seconds, to serve sudden spike in concurrent requests with consistently fast performance




For every 24 hours that your main cluster is in use, you accrue a one-hour credit for Concurrency Scaling. This means that Concurrency Scaling is free for >97% of customers.

How it works:

- 1 All queries go to the leader node, user only sees less wait for queries
- 2 When queries in designated WLM queue begin queuing, Amazon Redshift automatically routes them to the new clusters, enabling Concurrency Scaling automatically
- 3 Amazon Redshift automatically spins up a new cluster, processes waiting queries, and automatically shuts down the Concurrency Scaling cluster

Concurrency scaling

Amazon Redshift automatically adds transient clusters, in seconds, to serve sudden spike in concurrent requests with consistently fast performance




For every 24 hours that your main cluster is in use, you accrue a one-hour credit for Concurrency Scaling. This means that Concurrency Scaling is free for >97% of customers.

How it works:

- 1 All queries go to the leader node, user only sees less wait for queries
- 2 When queries in designated WLM queue begin queuing, Amazon Redshift automatically routes them to the new clusters, enabling Concurrency Scaling automatically
- 3 Amazon Redshift automatically spins up a new cluster, processes waiting queries, and automatically shuts down the Concurrency Scaling cluster

Concurrency scaling

Amazon Redshift automatically adds transient clusters, in seconds, to serve sudden spike in concurrent requests with consistently fast performance




For every 24 hours that your main cluster is in use, you accrue a one-hour credit for Concurrency Scaling. This means that Concurrency Scaling is free for >97% of customers.

How it works:

- 1 All queries go to the leader node, user only sees less wait for queries
- 2 When queries in designated WLM queue begin queuing, Amazon Redshift automatically routes them to the new clusters, enabling Concurrency Scaling automatically
- 3 Amazon Redshift automatically spins up a new cluster, processes waiting queries, and automatically shuts down the Concurrency Scaling cluster

Concurrency scaling

Amazon Redshift automatically adds transient clusters, in seconds, to serve sudden spike in concurrent requests with consistently fast performance



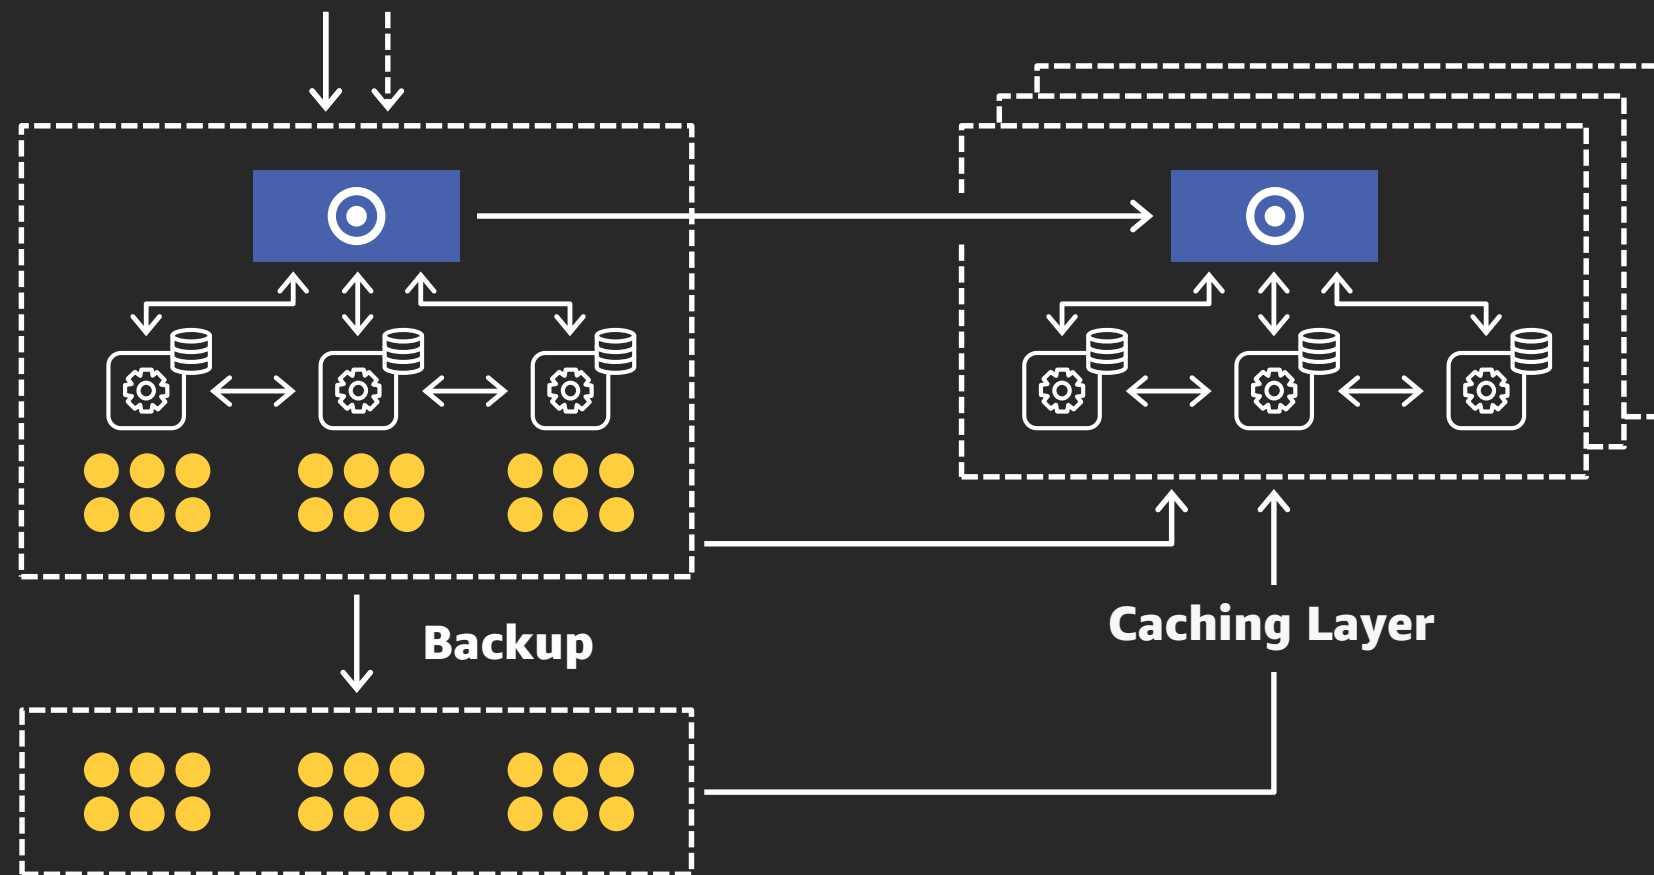
For every 24 hours that your main cluster is in use, you accrue a one-hour credit for Concurrency Scaling. This means that Concurrency Scaling is free for >97% of customers.

How it works:

- 1 All queries go to the leader node, user only sees less wait for queries
- 2 When queries in designated WLM queue begin queuing, Amazon Redshift automatically routes them to the new clusters, enabling Concurrency Scaling automatically
- 3 Amazon Redshift automatically spins up a new cluster, processes waiting queries, and automatically shuts down the Concurrency Scaling cluster

Concurrency scaling

Amazon Redshift automatically adds transient clusters, in seconds, to serve sudden spike in concurrent requests with consistently fast performance



For every 24 hours that your main cluster is in use, you accrue a one-hour credit for Concurrency Scaling. This means that Concurrency Scaling is free for >97% of customers.

How it works:

- 1 All queries go to the leader node, user only sees less wait for queries
- 2 When queries in designated WLM queue begin queuing, Amazon Redshift automatically routes them to the new clusters, enabling Concurrency Scaling automatically
- 3 Amazon Redshift automatically spins up a new cluster, processes waiting queries, and automatically shuts down the Concurrency Scaling cluster

Workload management: Example

Use case:

- Light ingestion/ELT on a continuous cadence of 10 minutes
- Peak reporting workload during business hours (7 am–7 pm)
- Heavy ingestion/ELT nightly (11 pm–3 am)

User types:

- Business reporting and dashboards
- Analysts and data science teams
- Database administrators

Workload management: Example manual WLM

- Enable: Short Query Acceleration
- Hidden superuser queue can be used by admins, manually switched into:
SET query_group TO 'superuser'
- The superuser queue has a single slot, the equivalent of 5–7% memory allocation, and no timeout

Create a queue for each workload type:

The screenshot displays the Databricks WLM configuration page. At the top, there is a checkbox for "Enable Short Query Acceleration for queries whose maximum runtime is" set to "dynamic". Below this, it shows "Max Concurrency Scaling clusters: 1". The interface lists three queues:

- Queue 1:** Memory (%) is 20, Concurrency on main is 2, Concurrency Scaling mode is Off, Timeout (ms) is empty. User groups include "Ingestion". Query groups are empty.
- Queue 2:** Memory (%) is 50, Concurrency on main is 10, Concurrency Scaling mode is Auto, Timeout (ms) is 120,000. User groups include "Dashboard". Query groups are empty.
- Default queue:** Memory (%) is 30, Concurrency on main is 3, Concurrency Scaling mode is Auto, Timeout (ms) is empty. User groups are empty. Query groups are empty.

Each queue section includes a "Delete" button, expand/collapse arrows, and a "Query Monitoring Rules" section with "Add rule from templates" and "Add custom rule" options.

Terminology and concept: Dynamic WLM

Manual WLM dynamic attributes

- Percent of memory

- Concurrency/queue slots

- Concurrency scaling

- Query timeout

- Enable short query acceleration

Changes to dynamic properties does not require a restart, it's a simple API call

Dynamic Workload Management Utility

<https://github.com/awslabs/amazon-redshift-utils/tree/master/src/WorkloadManagementScheduler>

WLM: Example (11 pm–3 am)

Enable: Short Query Acceleration

Increase memory and concurrency for ingestion queue

Decrease memory and concurrency for dashboard and default queues

☒ Enable Short Query Acceleration for queries whose maximum runtime is dynamic [Learn more](#)

Max Concurrency Scaling clusters: 1 [Edit](#)

Queue 1 Delete ^ v

Memory (%) ⓘ 80	Concurrency on main 5	Concurrency Scaling mode ⓘ Off v	Timeout (ms) ⓘ 	User groups <input type="checkbox"/> Match wildcards	Query groups <input type="checkbox"/> Match wildcards
				Ingestion +	+

▼ Query Monitoring Rules (0) [Add rule from templates](#) | [Add custom rule](#)

No rules have been defined.

Queue 2 Delete ^ v

Memory (%) ⓘ 10	Concurrency on main 1	Concurrency Scaling mode ⓘ Auto v	Timeout (ms) ⓘ 12000	User groups <input type="checkbox"/> Match wildcards	Query groups <input type="checkbox"/> Match wildcards
				Dashboard +	+

▼ Query Monitoring Rules (0) [Add rule from templates](#) | [Add custom rule](#)

No rules have been defined.

Default queue Delete ^ v

Memory (%) ⓘ 10	Concurrency on main 1	Concurrency Scaling mode ⓘ Off v	Timeout (ms) ⓘ
--------------------	--------------------------	--	--------------------

Automatic workload management (Auto WLM)

Allows for prioritization of different query workload

Goals

- Simplify WLM

Automatically controls concurrent number of executing queries

Automatically divides cluster memory

Auto WLM: Example

☒ Enable **Short Query Acceleration** for queries whose maximum runtime is dynamic [Learn more](#)

Max Concurrency Scaling clusters: 1 [Edit](#)

Default queue				Delete	^	v
Priority	Memory (%) i	Concurrency on main	Concurrency Scaling mode i			
Normal ^ v	Auto	Auto	Auto ^ v			
▼ Query Monitoring Rules (0)				Add rule from templates	Add custom rule	
No rules have been defined.						

Automatically manages memory allocation and concurrency of queries

Query monitoring rules (QMR)

Extension of workload management (WLM)

Allow the automatic handling of runaway (poorly written) queries

Rules applied to a WLM queue allow queries to be:

- LOGGED

- ABORTED

- HOPPED

Goals

- Protect against wasteful use of the cluster

- Log resource-intensive queries

Query monitoring rules (QMR)

Metrics with operators and values (e.g., return_row_count > 100000000) create a predicate

Multiple predicates can be AND-ed together to create a rule

Multiple rules can be defined for a queue in WLM. These rules are OR-ed together

```
If { rule } then [action]
{ rule : metric operator value } e.g.: rows_scanned > 10000000
    Metric:  cpu_time, query_blocks_read, rows scanned, query execution time, cpu &
             io skew per slice, join_row_count, etc.
    Operator:  <, >, ==
    Value:    integer
[action]: hop, log, abort
```

Best practices: WLM and QMR

Use Auto WLM—if you aren't sure how to set up WLM or your workload is highly unpredictable, or you are using the old default WLM

Use manual WLM—if you understand your workload patterns or require throttling certain types of queries depending on the time of day

Keep the number of WLM queues to a minimum, typically just three queues to avoid having unused queues

https://github.com/awslabs/amazon-redshift-utils/blob/master/src/AdminScripts/wlm_apex_hourly.sql

Use WLM to limit ingestion/ELT concurrency to two to three

To maximize query throughput, use WLM to throttle the number of concurrent queries to 15 or less

Use QMR rather than WLM to set query timeouts

Use QMR to log long-running queries

Save the superuser queue for administration tasks and canceling queries

Cluster sizing and resizing

Sizing Amazon Redshift cluster for production

Estimate the uncompressed size of the incoming data

Assume 3x compression (actual can be >4x)

Target 30–40% free space (resize to add/remove storage as needed)

Disk utilization should be at least 15% and less than 80%

Based on performance requirements, pick SSD or HDD

If required, additional nodes can be added for increased performance

Example: 20 TB of uncompressed data \approx 6.67 TB compressed

Depending on performance requirements, recommendation:

2-6xRA3.4xlarge or 4xDC2.8xlarge or 5xDS2.xlarge \approx 10TB of capacity

Resizing Amazon Redshift

Classic resize

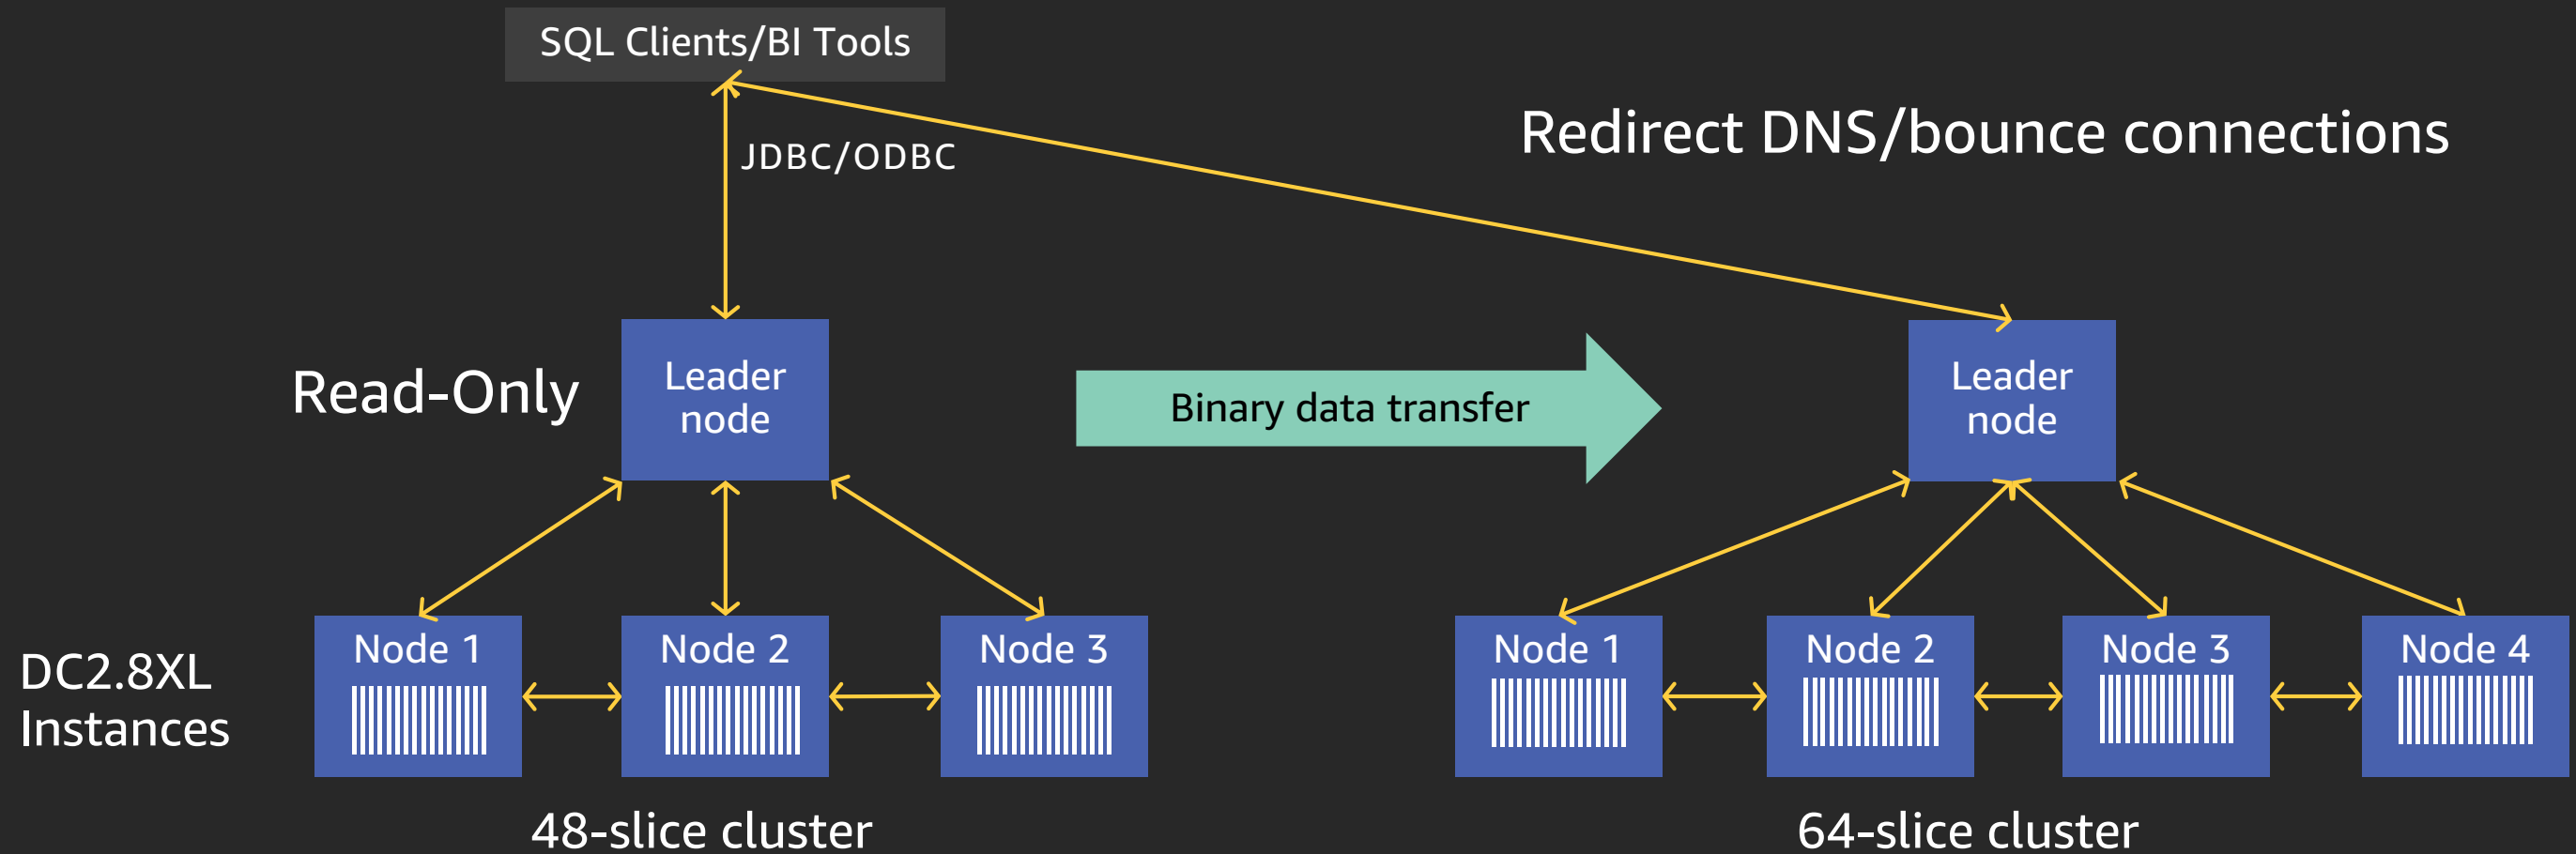
- Data is transferred from old cluster to new cluster (within hours)

- Change node types

Elastic resize

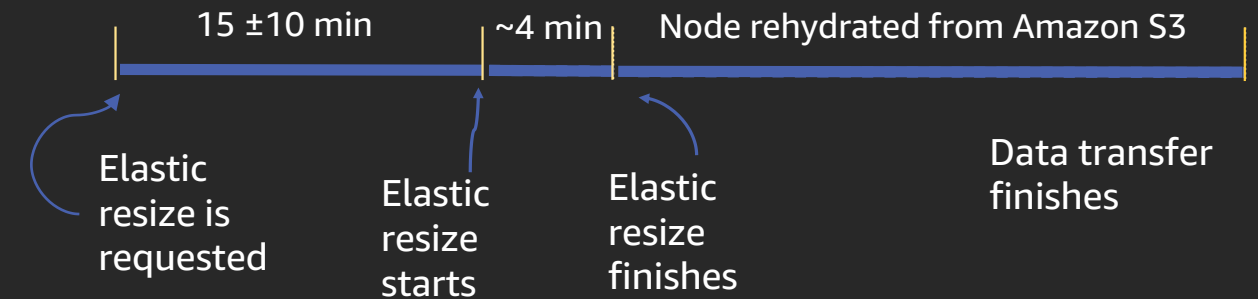
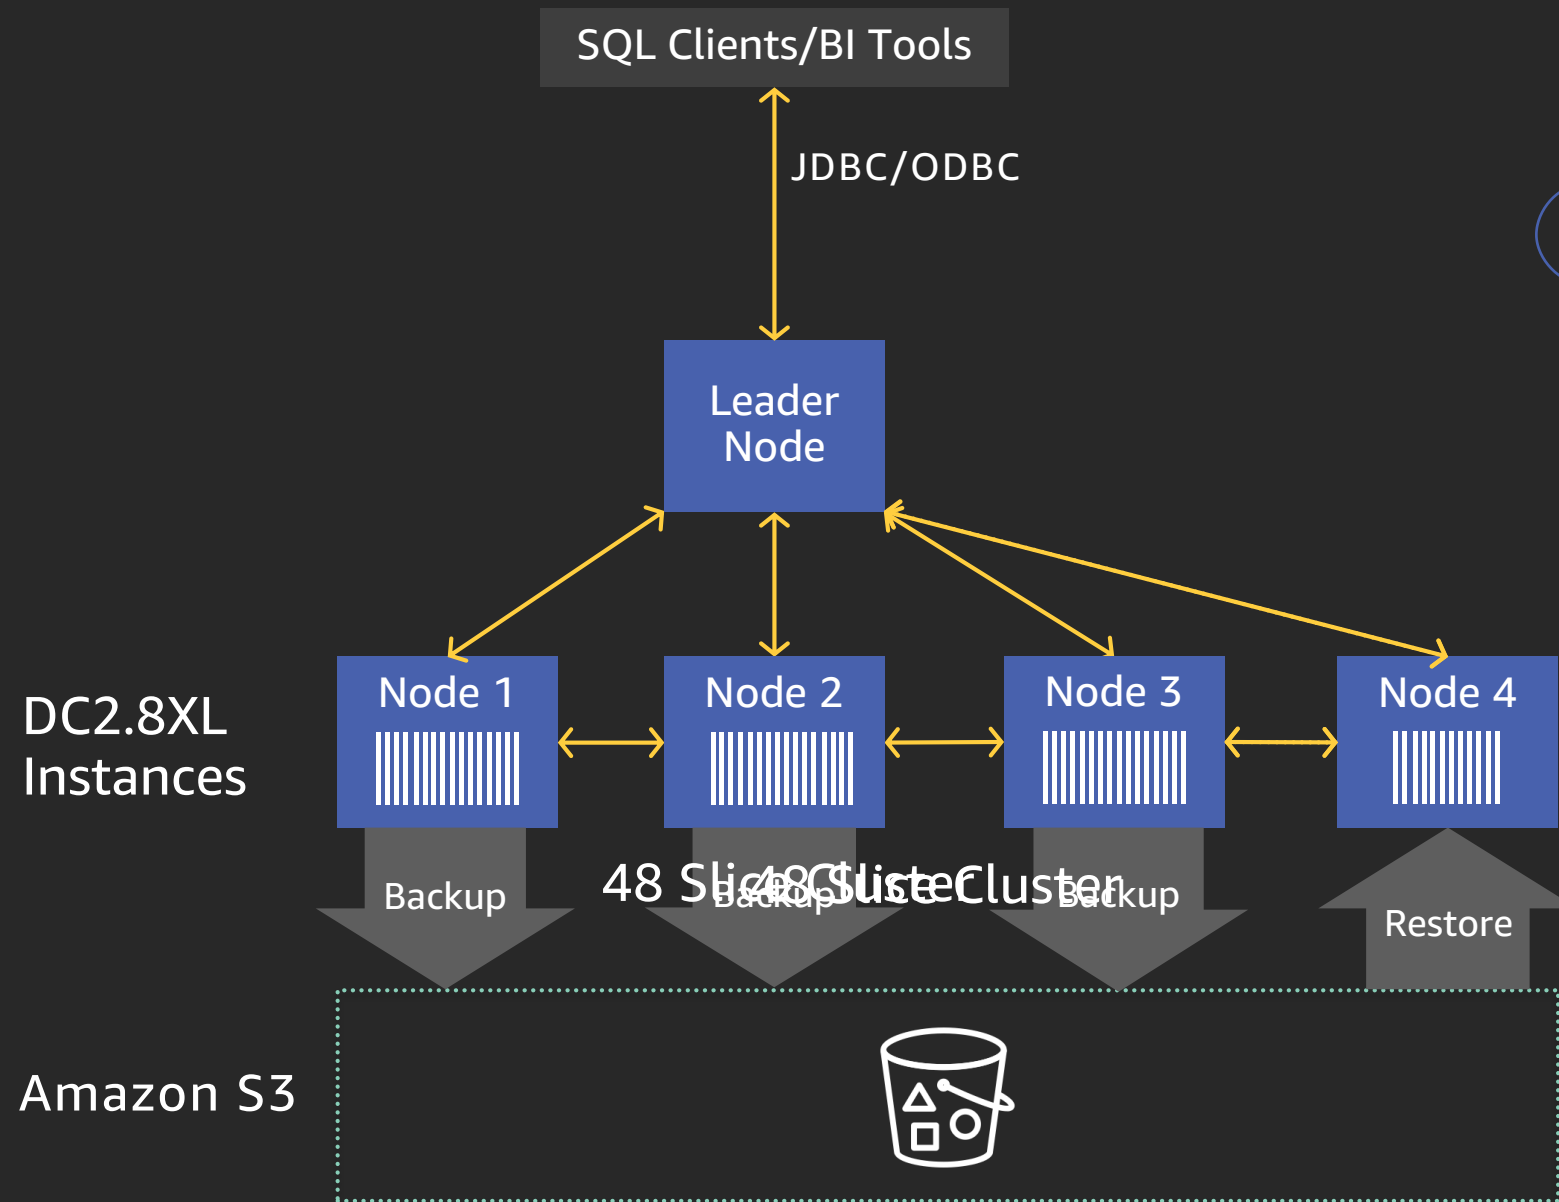
- Nodes are added/removed to/from existing cluster (within minutes)

Classic resize



- Source cluster is placed into read-only mode during resize
- All data is copied and redistributed on the target cluster
- Allows for changing node types

Elastic resize



- Cluster is fully available and ready to take traffic during the resize process
- No downtime or data loss
- S3 and provision the new node(s)
- No downtime or data loss
- Cluster is fully available for reads and writes
- Some queries within transactions maybe rollback

Elastic resize node increments

Instance type		Allowed increments	Max change from original size	Example: valid sizes for 4-node cluster
RA3 4xlarge	DC2 large	2x or 1/2 original cluster size only	Double, 1/2 size	2, 4, 8
DS2 xlarge				
RA3 16xlarge	DC2 8xlarge	Can allow \pm single node increments so long as slices remain balanced	Double, 1/2 size	2, 3, 4, 5, 6, 7, 8
DS2 8xlarge				

When to use elastic vs. classic resize

	Elastic resize	Classic resize
Scale up and down for workload spikes	✓	
Incrementally add/remove storage	✓	
Change cluster instance type (SSD \longleftrightarrow HDD)		✓
If elastic resize is not an option because of sizing limits		✓
Limited availability during resize	<5 minutes (parked connections)	1–24 hours (read-only)

Best practices: Cluster sizing

Use at least two compute nodes (multi-node cluster) in production for data mirroring

- Leader node is given for no additional cost

Maintain at least 20% free space or three times the size of the largest table

- Scratch space for usage, rewriting tables
- Free space is required for vacuum to re-sort table
- Temporary tables used for intermediate query results

The maximum number of available Amazon Redshift Spectrum nodes is a function of the number of slices in the Amazon Redshift cluster

If you're using DS2 instances, migrate to RA3

If you're using DC1 instances, upgrade to the DC2 instance type

- Same price as DC1, significantly faster
- Reserved Instances can be migrated without additional cost in the AWS Console

Amazon Redshift Advisor

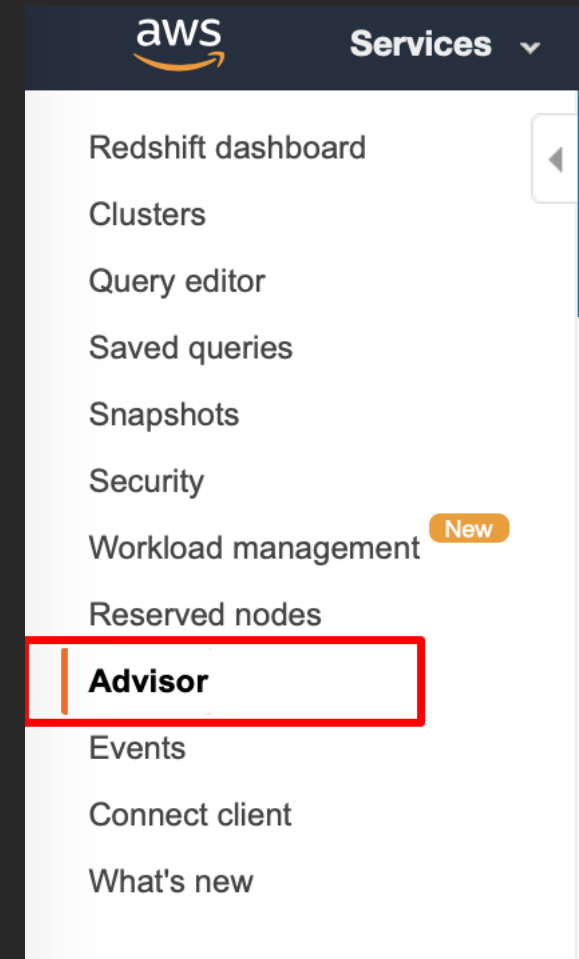
Amazon Redshift Advisor

Amazon Redshift Advisor available in Amazon Redshift Console

Runs daily scanning operational metadata

Observes with the lens of best practices

Provides tailored high-impact recommendations to optimize your Amazon Redshift cluster for performance and cost savings



Amazon Redshift Advisor: Recommendations

Recommendations include

- Skip compression analysis during COPY
- Split Amazon S3 objects loaded by COPY
- Compress Amazon S3 file objects loaded by COPY
- Compress table data
- Reallocate Workload Management (WLM) memory
- Cost savings
- Enable short query acceleration
- Alter distribution keys on tables

▼ Improve Query Performance with Distribution Keys

Refreshed 65 hours ago

Checks for appropriate distribution keys on tables.

Significantly improve query performance by using [ALTER TABLE](#) to redistribute the tables identified in this recommendation.

Amazon Redshift distributes table rows throughout the cluster according to the table [distribution style](#). Tables with `KEY` distribution require a column as the distribution key (`DISTKEY`). The distribution of table rows are based on the `DISTKEY` column values.


An appropriate `DISTKEY` places a similar number of rows on each node and is frequently referenced in join conditions. An optimized join occurs when tables are joined on their `DISTKEY` columns, accelerating query performance.

Observation

An analysis of the cluster's workload between 2019-10-13 and 2019-12-02 (50 days), identified tables that will significantly benefit from a `KEY` distribution style.

Recommendation

Use the following blocks of SQL statements to redistribute tables with the recommended `DISTKEY` column. In order to realize a significant performance benefit, all SQL statements within a recommendation group must be implemented.

```
-- First redistribution group
-- Database: "dev"
ALTER TABLE /*dkru-e554b525-a39c-4973-b17d-5d479ccff796-g0-0*/  ALTER DISTSTYLE KEY L
```



Additional resources

AWS Labs on GitHub: Amazon Redshift

<https://github.com/awslabs/amazon-redshift-utils>

<https://github.com/awslabs/amazon-redshift-monitoring>

<https://github.com/awslabs/amazon-redshift-udfs>

Admin scripts

Collection of utilities for running diagnostics on your cluster

Admin views

Collection of utilities for managing your cluster, generating schema DDL, and so on

Analyze Vacuum utility

Utility that can be scheduled to vacuum and analyze the tables within your Amazon Redshift cluster

Column Encoding utility

Utility that will apply optimal column encoding to an established schema with data already loaded

AWS big data blog: Amazon Redshift

Amazon Redshift Engineering's Advanced Table Design Playbook

<https://aws.amazon.com/blogs/big-data/amazon-redshift-engineerings-advanced-table-design-playbook-preamble-prerequisites-and-prioritization/>

—Zach Christopherson

Top 10 Performance Tuning Techniques for Amazon Redshift

<https://aws.amazon.com/blogs/big-data/top-10-performance-tuning-techniques-for-amazon-redshift/>

—Ian Meyers and Zach Christopherson


Twelve Best Practices for Amazon Redshift Spectrum

<https://aws.amazon.com/blogs/big-data/10-best-practices-for-amazon-redshift-spectrum/>

—Po Hong and Peter Dalton

Learn databases with AWS Training and Certification

Resources created by the experts at AWS to help you build and validate database skills



25+ free digital training courses cover topics and services related to databases, including

- Amazon Aurora
- Amazon Neptune
- Amazon DocumentDB
- Amazon DynamoDB
- Amazon ElastiCache
- Amazon Redshift
- Amazon RDS



Validate expertise with the new AWS Certified Database—
Specialty beta exam

Visit aws.training

Thank you!



Please complete the session survey in the mobile app.