



## Apache Hadoop 3

Andrew Wang

[andrew.wang@cloudera.com](mailto:andrew.wang@cloudera.com)

Daniel Templeton

[daniel@cloudera.com](mailto:daniel@cloudera.com)

# Who We Are

---



Andrew Wang

- HDFS @ Cloudera
- Hadoop PMC Member
- Release Manager for Hadoop 3.0



Daniel Templeton

- YARN @ Cloudera
- Hadoop PMC Member

# An Abbreviated History of Hadoop Releases

Date	Release	Major Notes
2007-11-04	0.14.1	First release at the ASF
2011-12-27	1.0.0	Security, HBase support
2012-05-23	2.0.0	YARN, NameNode HA, wire compat
2014-11-18	2.6.0	HDFS encryption, rolling upgrade, node labels
2015-04-21	2.7.0	Truncate, Variable-length blocks, YARN Global Caching,
2017-03-22	2.8.0	Cloud improvement, Azure Data Lake, and etc.
2017-11-17	2.9.0	Stability Improvement
2017-12-13	3.0.0	Java 8, Erasure Coding, S3Guard, YARN Timeline Service

# Motivation for Hadoop 3

- Upgrade minimum Java version to Java 8
  - Java 7 end-of-life in April 2015
  - Many Java libraries now only support Java 8
- HDFS erasure coding
  - Major feature that refactored core pieces of HDFS
  - Too big to backport to 2.x
- Classpath isolation
  - Potentially impacts all clients
- Other miscellaneous incompatible bugfixes and improvements
  - Hadoop 2.x was branched in 2011
  - 6 years of changes waiting for 3.0

# Hadoop 3 Status and Release Plan

- After four alphas and one beta, 3.0.0 is out!
- Took close to two years from inception
- 3.0.1 and 3.1.0 are already in progress


Release	Date	
3.0.0-alpha1	2016-09-03	✓
3.0.0-alpha2	2017-01-25	✓
3.0.0-alpha3	2017-05-26	✓
3.0.0-alpha4	2017-07-07	✓
3.0.0-beta1	2017-10-03	✓
3.0.0 GA	2017-12-13	✓
3.0.1	2017 Mar	

<https://cwiki.apache.org/confluence/display/HADOOP/Hadoop+3.0.0+release>

# HDFS & Hadoop Features


# 3x replication vs. Erasure coding

/foo.csv - 3 block file



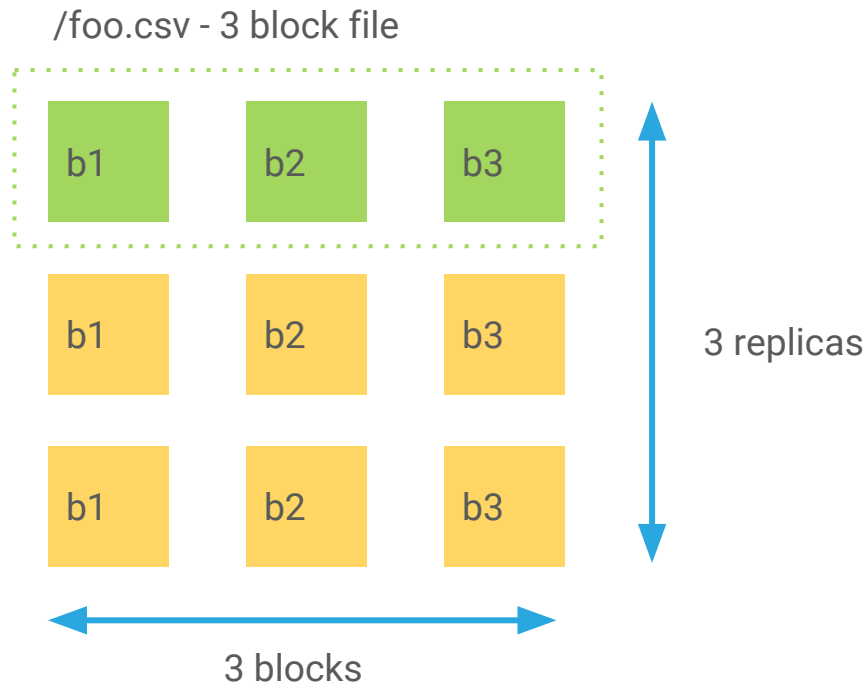
# 3x replication vs. Erasure coding

/foo.csv - 3 block file





# 3x replication vs. Erasure coding




$3 \times 3 = 9$  total replicas

$9 / 3 = 200\% \text{ overhead!}$


# 3x replication vs. Erasure coding

/foo.csv - 3 block file




# 3x replication vs. Erasure coding

/foo.csv - 3 block file



# 3x replication vs. Erasure coding



$$3 + 2 = 5 \text{ replicas}$$
$$5 / 3 = 67\% \text{ overhead!}$$

# 3x replication vs. Erasure coding

/foo.csv - 3 block file



3 data blocks




2 parity blocks

$3 + 2 = 5$  replicas  
 $5 / 3 = 67\%$  overhead!

/bigfoo.csv - 10 block file



10 data blocks






4 parity blocks

$10 + 4 = 14$  replicas  
 $14 / 10 = 40\%$  overhead!

# EC Reconstruction




/foo.csv - 3 block file



Reed-Solomon (3,2)

# EC Reconstruction


/foo.csv - 3 block file



Reed-Solomon (3,2)

# EC Reconstruction


/foo.csv - 3 block file



Reed-Solomon (3,2)

Read 3 remaining blocks

Run RS to recover b3



New copy of b3 recovered




## Erasure coding (HDFS-7285)

- Motivation: improve storage efficiency of HDFS
  - **~2x** the storage efficiency compared to 3x replication
  - Reduction of overhead from 200% to 40%
- Uses Reed-Solomon(k,m) erasure codes instead of replication
  - Support for multiple erasure coding policies
  - RS(3,2), RS(6,3), RS(10,4)
- Can improve data durability
  - RS(6,3) can tolerate 3 failures
  - RS(10,4) can tolerate 4 failures
- Missing blocks reconstructed from remaining blocks


## EC implications

- File data is striped across multiple nodes and racks
- Reads and writes are **remote** and **cross-rack**
- Reconstruction is **network-intensive**, reads  $m$  blocks cross-rack
- Important to use Intel's optimized ISA-L for performance
  - 1+ GB/s encode/decode speed, much faster than Java implementation
- Combine data into larger files to avoid an explosion in # replicas
  - Bad: 1x1GB file -> RS(10,4) -> 14x100MB EC blocks (4.6x # replicas)
  - Good: 10x1GB file -> RS(10,4) -> 14x1GB EC blocks (0.46x # replicas)
- Works best for archival / cold data use cases

# EC performance

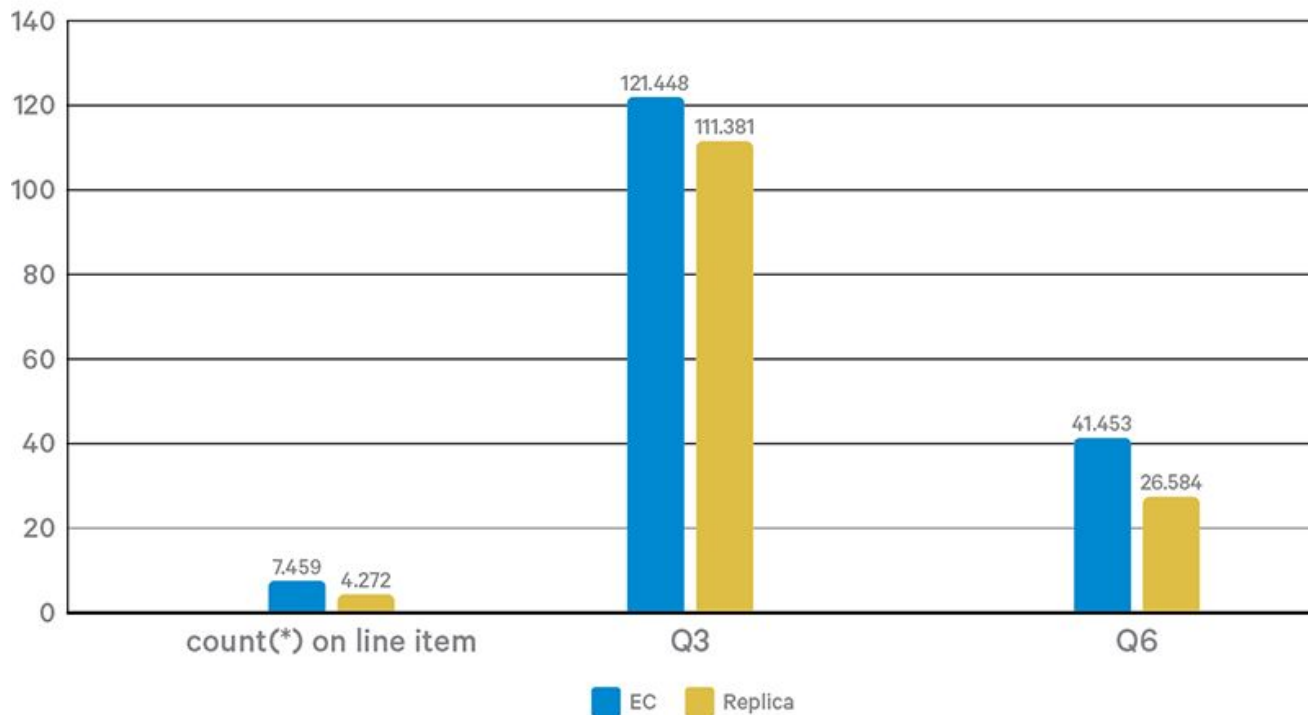


# EC performance



# EC performance

Spark TPC-H 500GB Query Time (s) — 2 DNs killed



# Erasure coding status

- Massive development effort by the Hadoop community
  - 20+ contributors from many companies
    - Cloudera, Intel, Hortonworks, Huawei, Y! JP, ...
  - 100s of commits over more than three years (started in 2014)
- Erasure coding is **ready in 3.0.0 GA!**
- Current focus is on testing and integration efforts
  - Want the complete Hadoop stack to work with HDFS erasure coding enabled
  - Ongoing stress / endurance testing to ensure stability at scale

# Classpath isolation (HADOOP-11656)

- Hadoop leaks lots of dependencies onto the application's classpath
  - Known offenders: Guava, Protobuf, Jackson, Jetty, ...
- No separate HDFS client jar means server jars are leaked
- YARN / MR clients not shaded
- **HDFS-6200**: Split HDFS client into separate JAR
- **HADOOP-11804**: Shaded hadoop-client dependency
- **YARN-6466**: Shade the task umbilical for a clean YARN container environment (ongoing)


# Miscellaneous

- Supportability improvements
  - Shell script rewrite
  - Intra-DataNode balancer
  - Move default ports out of the ephemeral range
- Support for multiple Standby NameNodes
- Cloud enhancements
  - Support for Microsoft Azure Data Lake and Aliyun OSS
  - S3 consistency and performance improvements
- Tightened Hadoop compatibility policy




# YARN Features

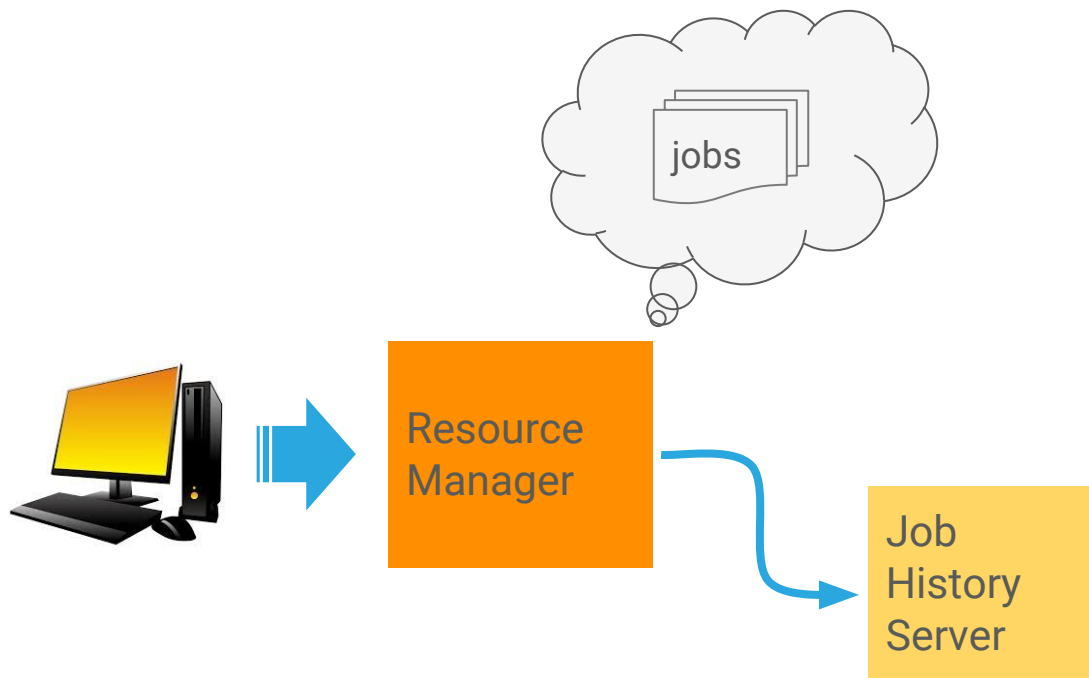
# Job History Server



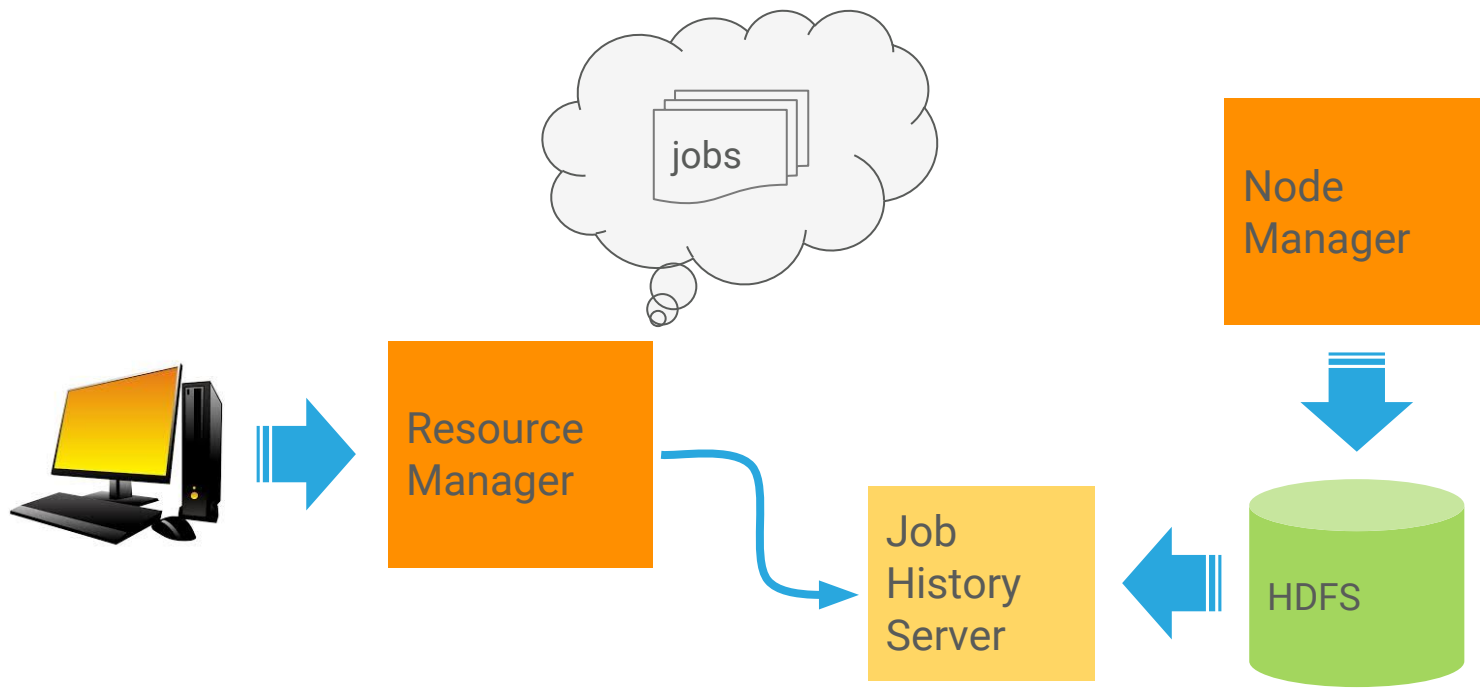
# Job History Server



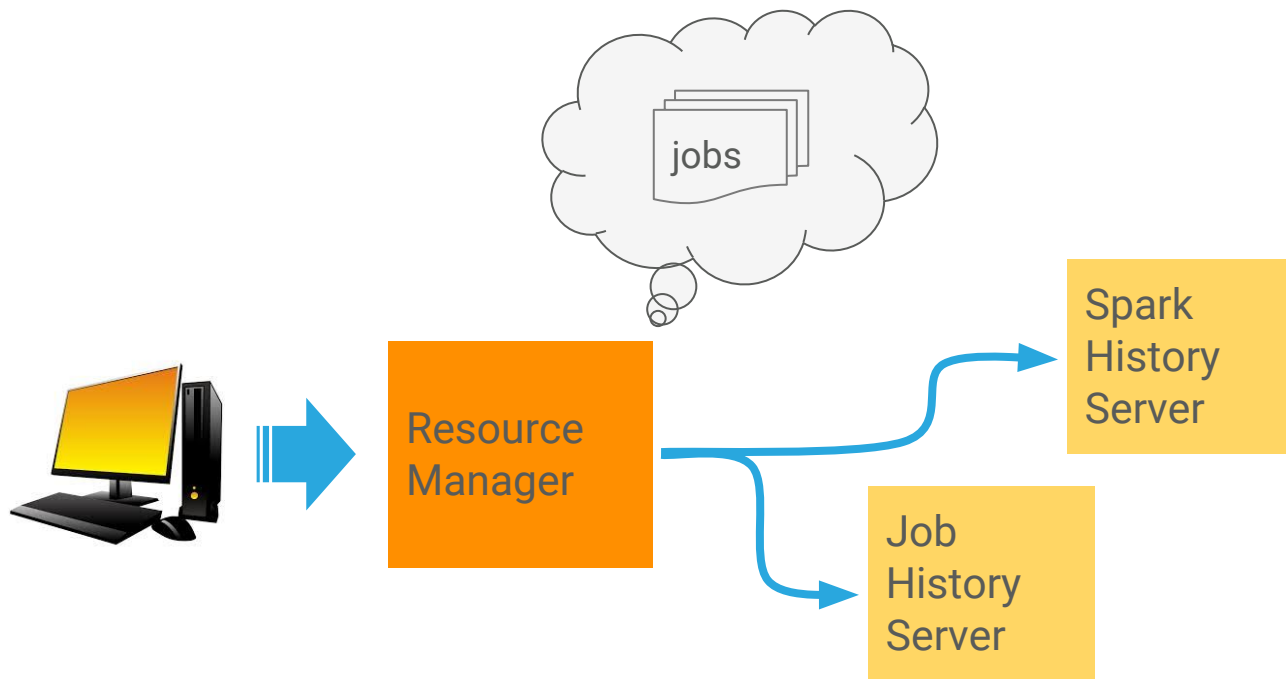
# Job History Server




# Job History Server



# Job History Server



# Job History Server

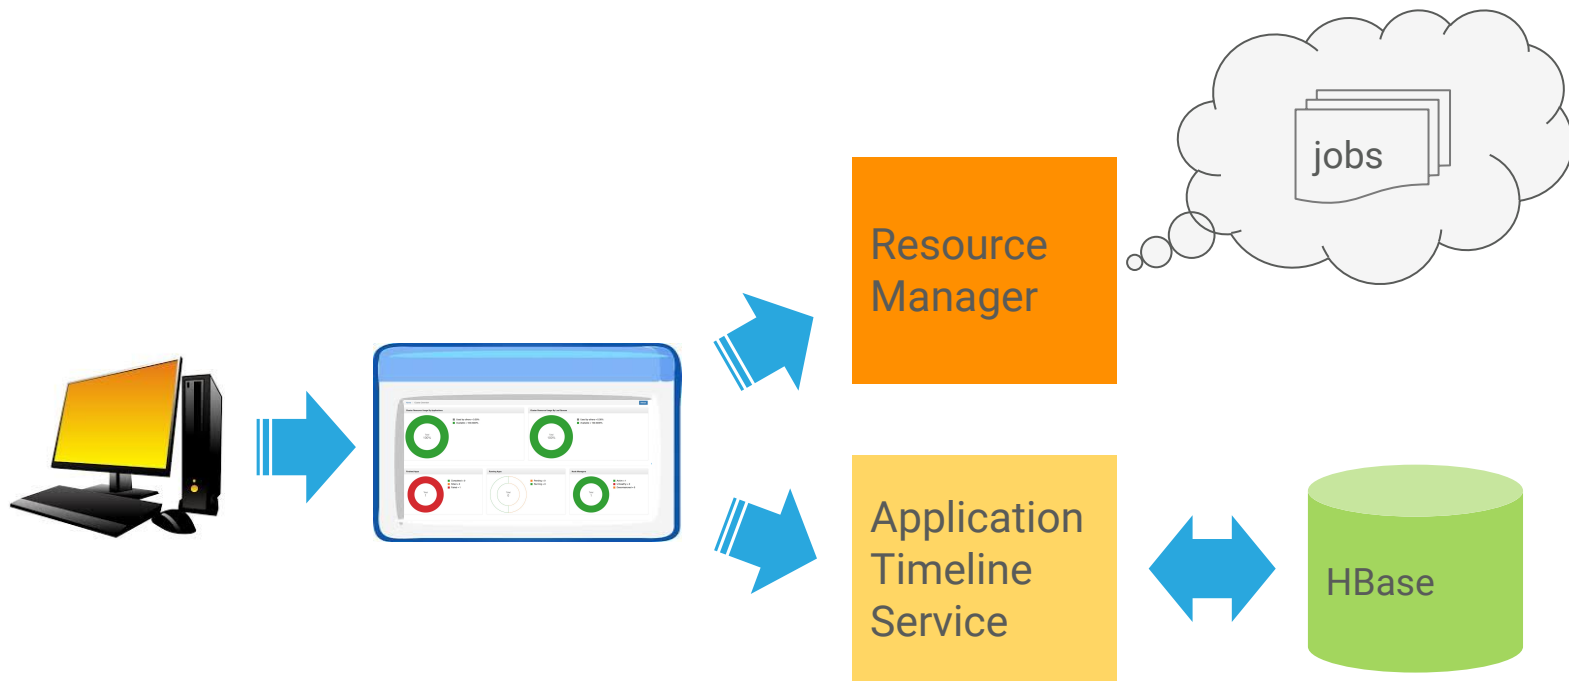


# Application Timeline Service v2

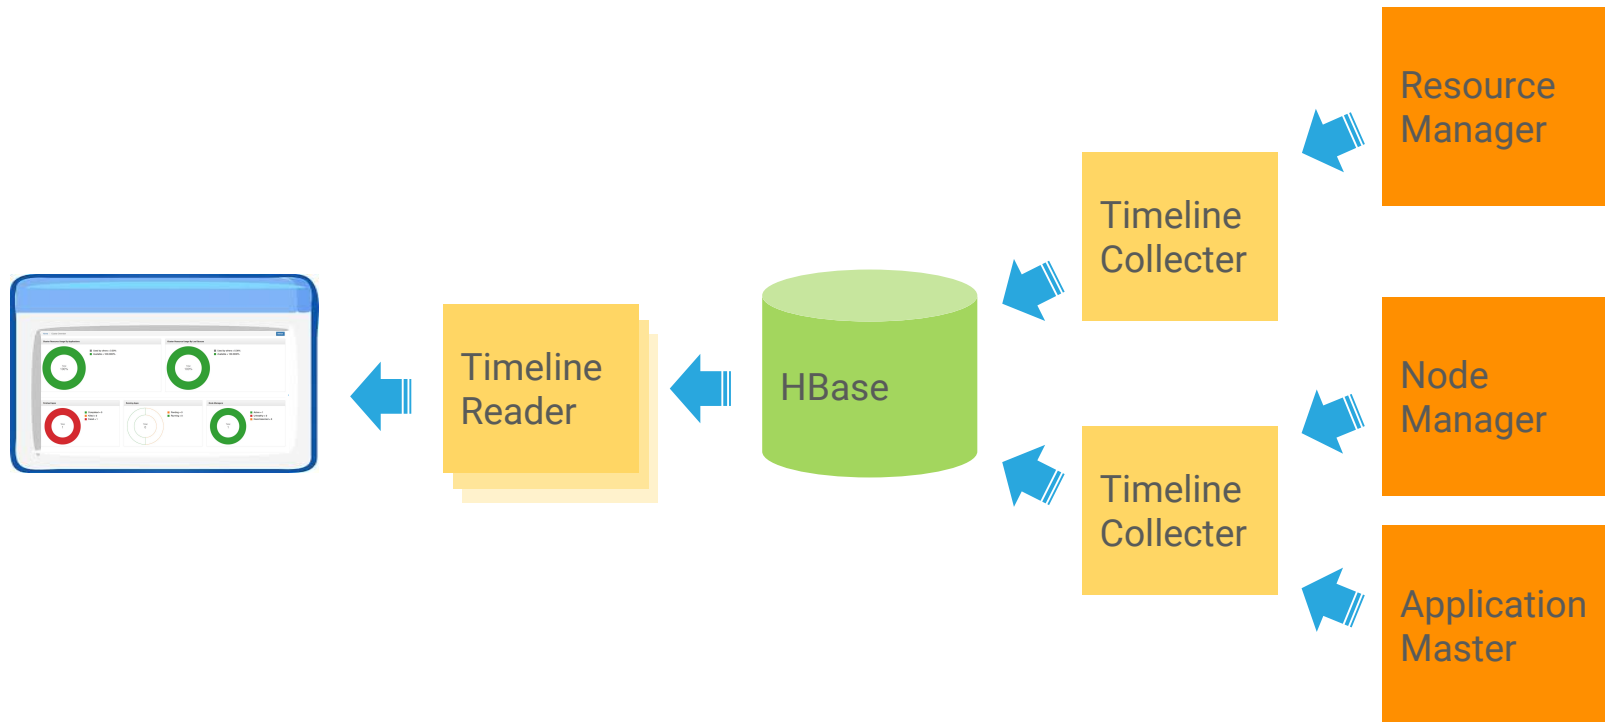
- Store for application and system events and data
  - Distributed
  - Scalable
  - Structured Data Model
- Updated in real time
  - Application status
  - Application metrics
  - System metrics
- Fed by resource manager, node manager, and application masters
- REST API



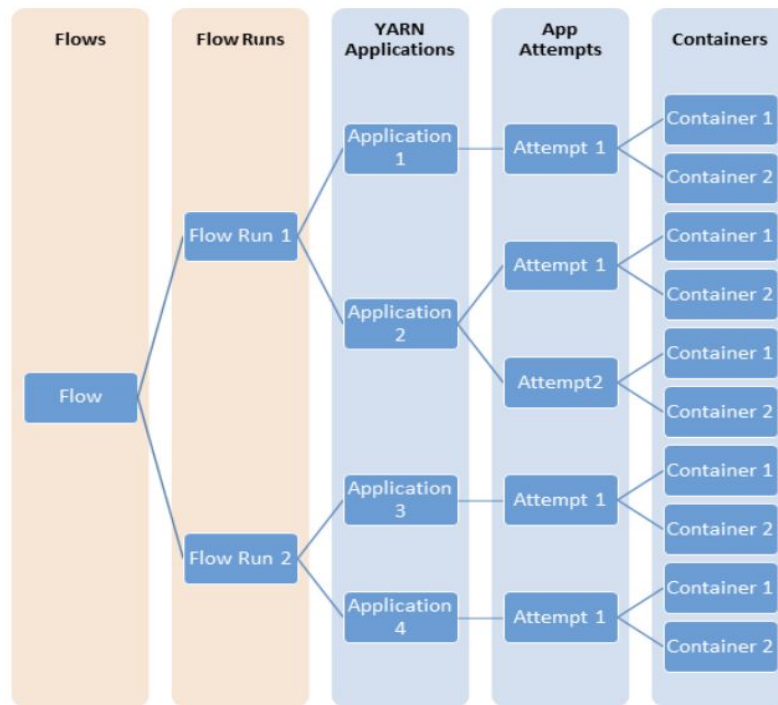
# Application Timeline Service v2



# Application Timeline Service v2



# Application Timeline Service v2 Flows



# Application Timeline Service v2 Flows

[Cluster Overview](#)[Queues](#)[Applications](#)[Services](#)[Flow Activity](#)[Nodes](#)[Home](#) / [Flow Activities](#)[Refresh](#)

Flow Activity

[Flow Activities](#)

## Recent Flow Activities

Flow Name	User	Flow ID	Last Execution Date
QuasiMonteCarlo	haibochen	haibochen-pseudo-tms-cluster/haibochen/QuasiMont...	2017/06/06
Sleep job	haibochen	haibochen-pseudo-tms-cluster/haibochen/Sleep job	2017/06/06
20170606_demo_1	haibochen	haibochen-pseudo-tms-cluster/haibochen/20170606...	2017/06/05
long sleep	haibochen	haibochen-pseudo-tms-cluster/haibochen/long sleep	2017/06/05
pi_calculation	haibochen	haibochen-pseudo-tms-cluster/haibochen/pi_calculat...	2017/06/05
word count	haibochen	haibochen-pseudo-tms-cluster/haibochen/word count	2017/06/05
TeraGen	haibochen	haibochen-pseudo-tms-cluster/haibochen/TeraGen	2017/05/25
mapreduce.BaileyBorweinPlouffe_10_20	haibochen	haibochen-pseudo-tms-cluster/haibochen/mapreduc...	2017/05/24
tsv2_demo_1	haibochen	haibochen-pseudo-tms-cluster/haibochen/tsv2_dem...	2017/05/24
tsv2_demo_2	haibochen	haibochen-pseudo-tms-cluster/haibochen/tsv2_dem...	2017/05/24

# Application Timeline Service v2 Flows



[Cluster Overview](#) [Queues](#) [Applications](#) [Services](#) [Flow Activity](#) [Nodes](#)

[Home](#) / [Flow Activities](#) / [Flow Info \[haibochen-pseudo-tms-cluster/haibochen/QuasiMonteCarlo\]](#)

[Refresh](#)

## Flow Information

[Flow Info](#)

[Flow Runs](#)

## Flow Runs

[Add Metrics Widget](#)

Search...


Search

1

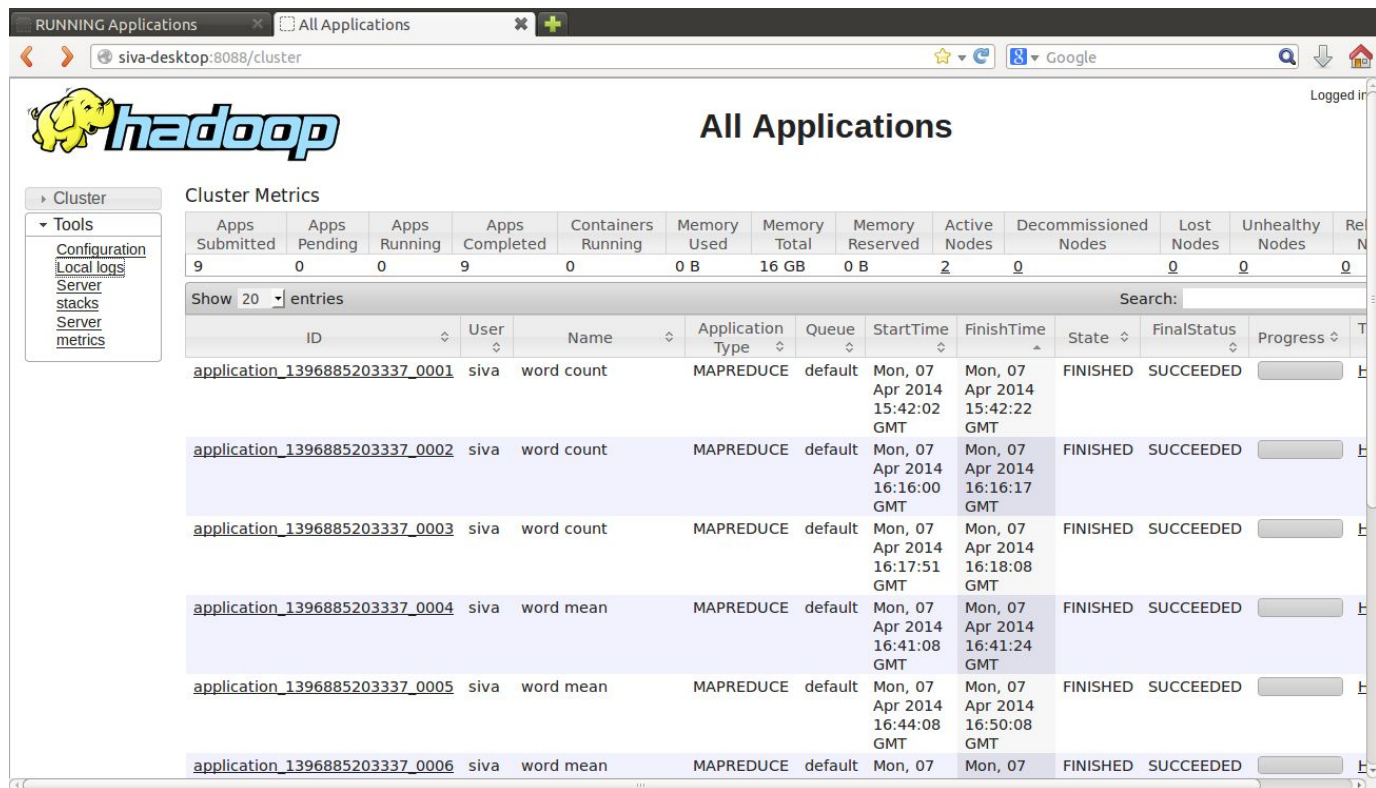
10 Rows

Run ID	Run Duration	CPU VCores	Memory Used	Creation Time	End Time
haibochen@QuasiMonteCarlo/1496868566...	31 Secs	N/A	N/A	2017/06/07 13:49:26	2017/06/07 13:49:57
haibochen@QuasiMonteCarlo/1495747118...	32 Secs	N/A	N/A	2017/05/25 14:18:38	2017/05/25 14:19:10
haibochen@QuasiMonteCarlo/1495733415...	28 Secs	N/A	N/A	2017/05/25 10:30:15	2017/05/25 10:30:44

Flow Run Vs Run Duration



# Old YARN UI



The screenshot shows the 'All Applications' page in the Old YARN UI. The page features a sidebar with navigation links, a header with the Hadoop logo, and a main content area with cluster metrics and a table of applications.

**Cluster Metrics**

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rel N
9	0	0	9	0	0 B	16 GB	0 B	2	0	0	0	0

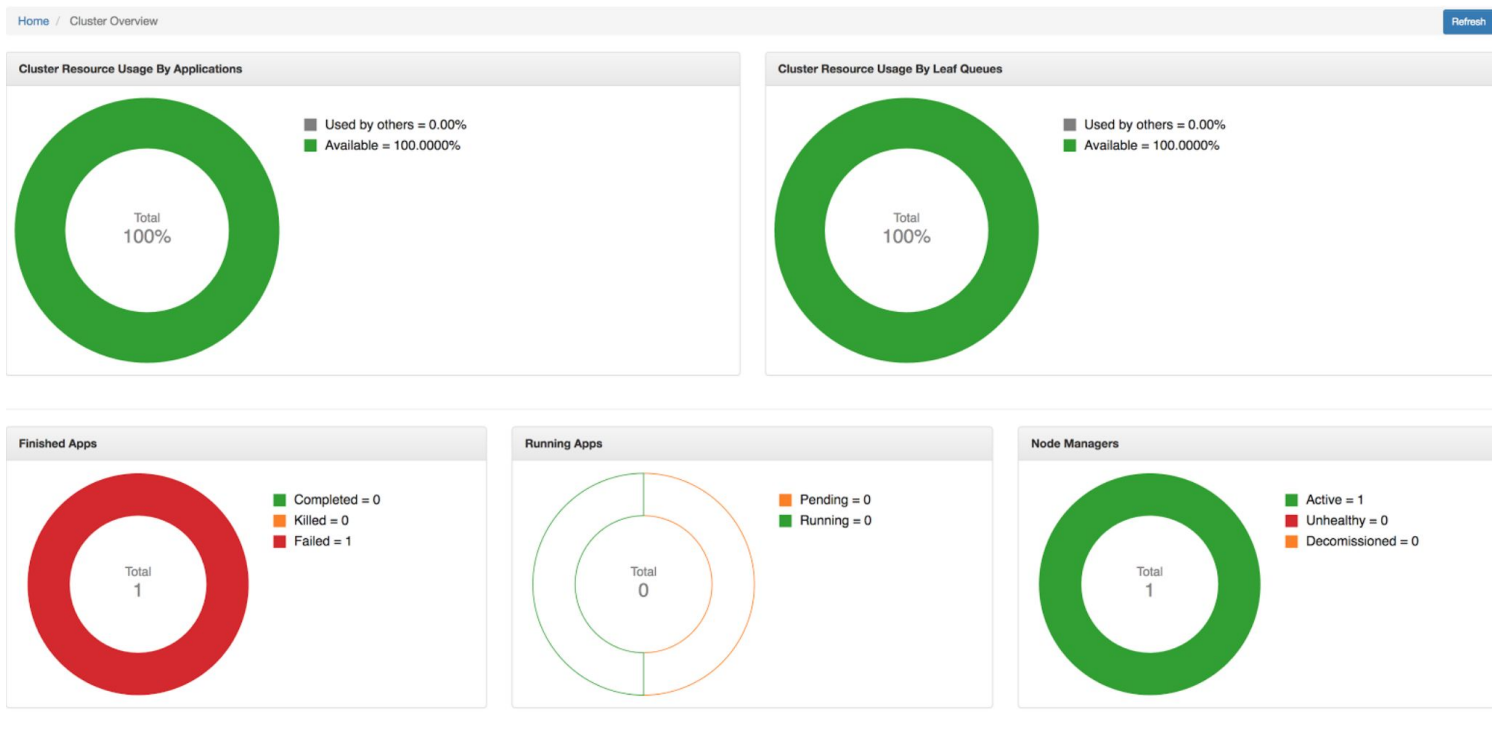
**Applications Table**

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress
<a href="#">application_1396885203337_0001</a>	siva	word count	MAPREDUCE	default	Mon, 07 Apr 2014 15:42:02 GMT	Mon, 07 Apr 2014 15:42:22 GMT	FINISHED	SUCCEEDED	<div></div>
<a href="#">application_1396885203337_0002</a>	siva	word count	MAPREDUCE	default	Mon, 07 Apr 2014 16:16:00 GMT	Mon, 07 Apr 2014 16:16:17 GMT	FINISHED	SUCCEEDED	<div></div>
<a href="#">application_1396885203337_0003</a>	siva	word count	MAPREDUCE	default	Mon, 07 Apr 2014 16:17:51 GMT	Mon, 07 Apr 2014 16:18:08 GMT	FINISHED	SUCCEEDED	<div></div>
<a href="#">application_1396885203337_0004</a>	siva	word mean	MAPREDUCE	default	Mon, 07 Apr 2014 16:41:08 GMT	Mon, 07 Apr 2014 16:41:24 GMT	FINISHED	SUCCEEDED	<div></div>
<a href="#">application_1396885203337_0005</a>	siva	word mean	MAPREDUCE	default	Mon, 07 Apr 2014 16:44:08 GMT	Mon, 07 Apr 2014 16:50:08 GMT	FINISHED	SUCCEEDED	<div></div>
<a href="#">application_1396885203337_0006</a>	siva	word mean	MAPREDUCE	default	Mon, 07	Mon, 07	FINISHED	SUCCEEDED	<div></div>

# New YARN UI

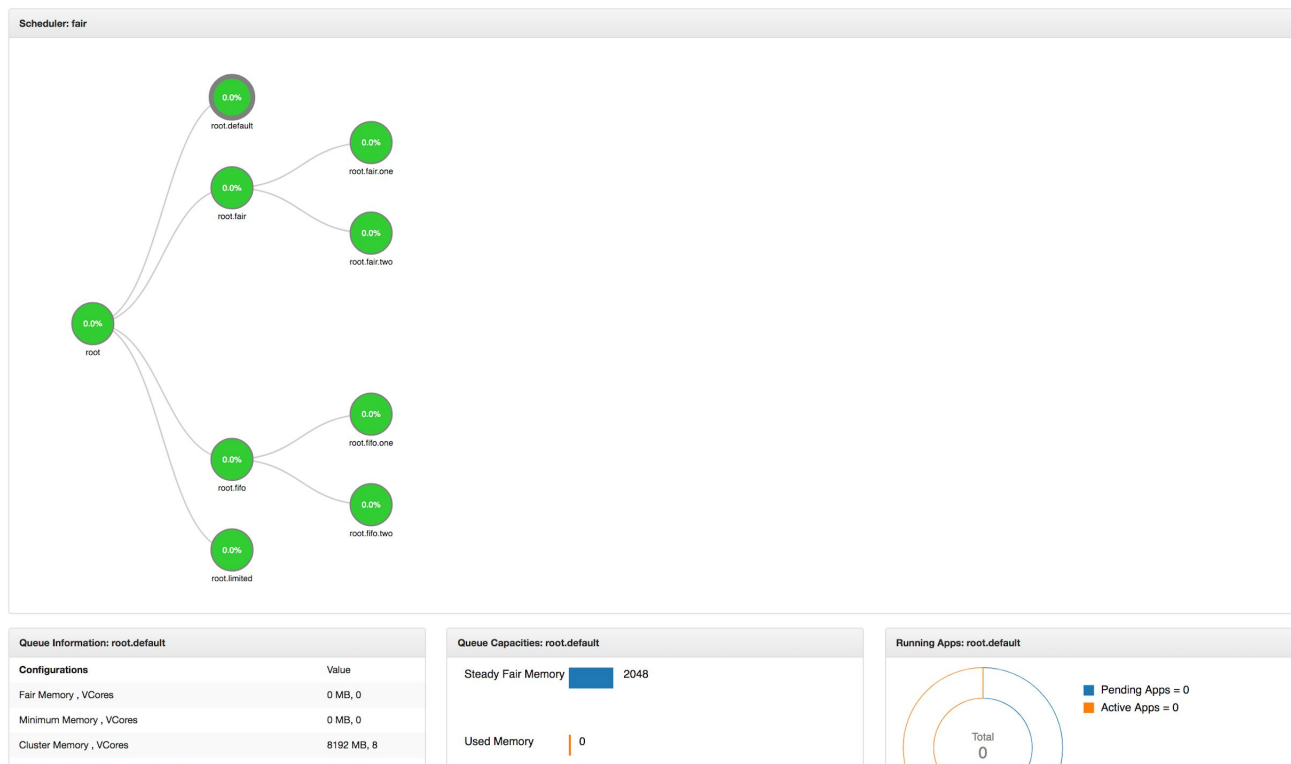
- Rich client application
  - Built on Node.js and Ember
- Improved visibility into cluster usage
  - Memory, CPU
  - By queues and applications
  - Sunburst graphs for hierarchical queues
  - NodeManager heatmap
- ATSV2 integration
  - Plot container start/stop events
  - Easy to capture delays in app execution

# New YARN UI: Cluster Overview





# New YARN UI: Queues



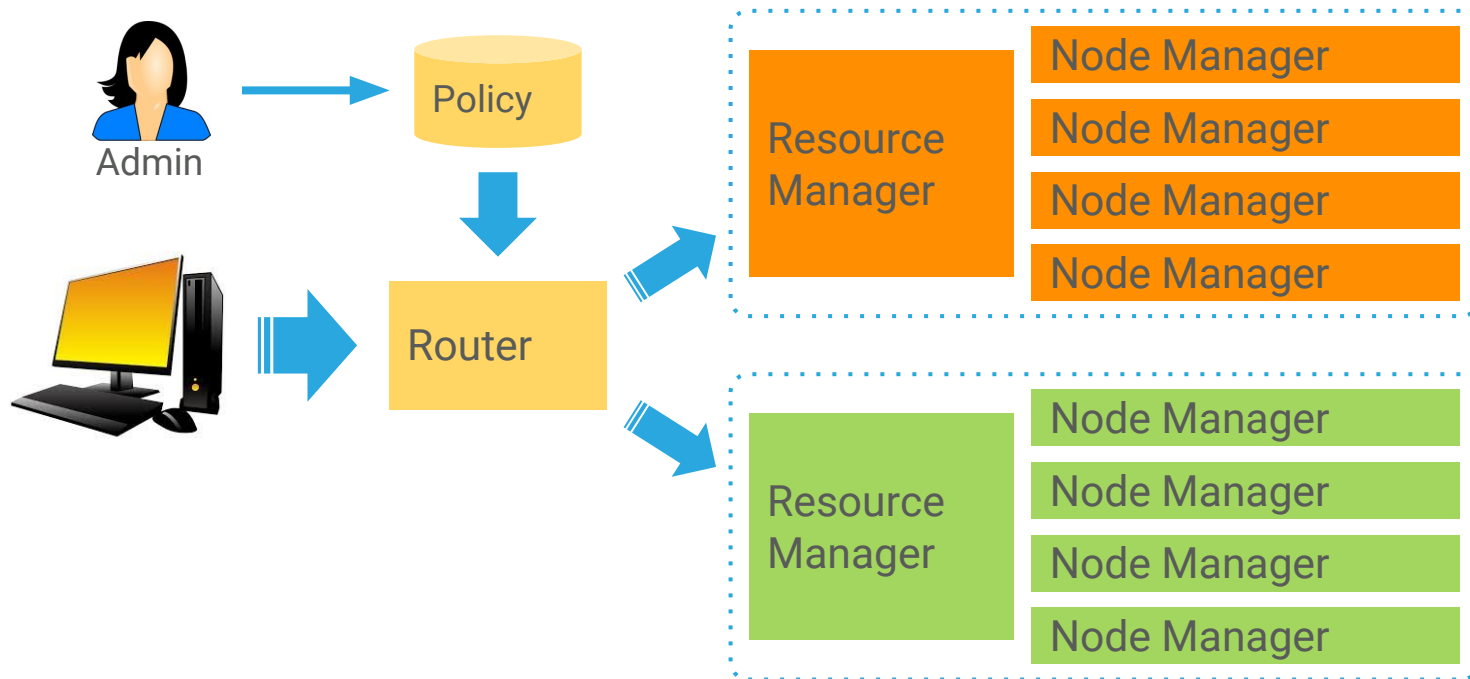
# Resource Types

- Before Hadoop 3 memory and CPU are the only managed resources
- Resource Types allows adding new managed resources
  - Countable resources: GPUs, Disks etc.
  - Static resources: Java version, Python version, hardware profile, ...
    - *Still in proposal stage*
- Resource profiles
  - Similar conceptually to EC2 instance types
  - Capture complex resource request
- DRF for scheduling
- Current virtual CPU cores and memory resources work as before

# YARN Federation

- YARN scalability
  - Twitter runs a 10k node cluster with fair scheduler
  - Yahoo! runs 4k node cluster with capacity scheduler
- Federation
  - Restrict users to sub-clusters based on policy
  - Scalability to 100k nodes and beyond
  - Independent cluster scheduling

# YARN Federation




# Opportunistic Containers

- Scheduler's job is to keep all resources busy
- Scheduling gaps
  - Nothing to run
  - Resource contention
  - Resource reservations
- Opportunistic containers fill those gaps
  - Requested explicitly
  - Dedicated scheduler
  - Queued at the node managers
  - Scheduled locally when resources are available
  - Preempted when guaranteed containers need to run
- *Coming in 2.9 and 3.0*

# Oversubscription

- Resource utilization is typically low in most clusters (20-50%)
  - Provision for peak usage
- Usage < Allocation
  - Mean Usage =  $\frac{1}{2}$  Peak Usage



# Oversubscription

- Oversubscription
  - Allocate opportunistic containers to use *allocated-but-unused* resources
  - Jobs automatically use these unless they *opt-out*
  - Threshold to control aggressiveness of oversubscription
  - Threshold to trigger preemption
- *Currently in progress*

## Other YARN Improvements

- Long Running Services
  - Slider merging into YARN
  - Docker support
- Scheduler improvements
  - Capacity scheduler
    - Performance and preemption improvements
    - Online scheduling (“global scheduler”)
    - Queue management
  - Fair scheduler
    - Performance and preemption improvements
- High availability improvements
  - Better handling of transient network issues
  - ZK-store scalability: Limit number of children under a znode
- MapReduce Native Collector (MAPREDUCE-2841)
  - Native implementation of the map output collector
  - Up to 30% faster for shuffle-intensive jobs



# Summary: What's new in Hadoop 3.0?

- Storage Optimization
  - HDFS: Erasure codes
- Improved Visibility into Cluster Operations
  - YARN: ATSV2
  - YARN: New UI
- Scalability & Multi-tenancy
  - YARN: Federation
- Improved Utilization
  - YARN: Opportunistic Containers
  - YARN: Oversubscription
- Refactor Base
  - Lots of Trunk content
  - JDK8 and newer dependent libraries

# Compatibility and Testing

# Compatibility

- Strong feedback from large users on the need for compatibility
- Preserves **wire-compatibility** with Hadoop 2 clients
  - Impossible to coordinate upgrading **off-cluster Hadoop clients**
- Will support **rolling upgrade** from Hadoop 2 to Hadoop 3
  - Can't take **downtime** to upgrade a business-critical cluster
- Not fully preserving API compatibility!
  - Dependency version bumps
  - Removal of deprecated APIs and tools
  - Shell script rewrite, rework of Hadoop tools scripts
  - Incompatible bug fixes

# Testing and Validation

- Cloudera CDH 6 is based on upstream Hadoop 3.0.0
  - Running full test suite
  - Integration of Hadoop 3 with all components in CDH stack
  - Same integration tests used to validate CDH5
- Plans for extensive HDFS EC testing by Cloudera and Intel
- Happy synergy between 2.8.x and 3.0.x lines
  - Shares much of the same code, fixes flow into both
  - Yahoo! doing scale testing of 2.8.0

# Conclusion

- Hadoop 3.0.0 GA is out!
- Shiny new features
  - HDFS erasure coding
  - Client classpath isolation
  - YARN ATSV2
  - YARN Federation
  - Opportunistic containers and oversubscription
- Great time to get involved in testing and validation

# Thank you

---

Andrew Wang  
andrew.wang@cloudera.com

Daniel Templeton  
daniel@cloudera.com